

Tensor Omics (TOX)

Analyze gene expression data and evolutionary relationships between genes

User's Guide

Vivian Bass¹, Anna Beketova¹, Luka Faensen¹, Asis Hallab¹, Aaron Schröder¹, Franz-Eric Sill¹

¹University of Applied Sciences Bingen, Germany

Last revised November 11, 2025

Contents

1	Tensor Omics Method and Theory	3
1.1	From Raw Data to Expression Vectors	3
1.2	Normalization of Expression Values	3
1.3	Calculation of family Centroids	3
1.4	Euclidean Distance Between Expression Vectors	4
1.5	Outlier Detection Based on Family-Scaled Distances	4
1.6	Tissue Versatility and Specificity	5
2	Error Handling	5
3	Implementation Notes	6
3.1	Fortran Pipeline	7
3.1.1	Overview	7
3.1.2	Data Management	7
3.1.3	Data Persistence	13
3.1.4	Data Processing	25
3.1.5	Family-Based Analysis	28
3.1.6	Relative Axis Plane Projection	33
3.1.7	Contribution Analysis	33
3.1.8	Fortran Utils	39
3.2	Python Pipeline	41
3.3	R Pipeline	63
3.4	Fortran utilities	80

1 Tensor Omics Method and Theory

The Tensor Omics (TOX) tool implements a streamlined pipeline for analyzing gene expression data and evolutionary relationships between genes. The pipeline integrates tools for data normalization, and geometric analysis of gene expression vectors. This document outlines the current implementation details and provides the mathematical formulation of the algorithms used in the prototype.

1.1 From Raw Data to Expression Vectors

The starting point for all downstream analyses is a table of gene-wise expression measurements across a set of biological samples, such as tissues, developmental stages, or experimental conditions. These measurements may originate from transcriptomic, proteomic, or metabolomic experiments, typically normalized to account for sequencing depth and technical variation.

What is Gene Expression Data? For each gene, a real-valued quantity is recorded that reflects its biological activity in a given sample. In the case of transcriptomics, this is often the abundance of messenger RNA (mRNA) molecules transcribed from that gene, estimated via sequencing and mapped back to the gene’s known sequence. The result is a non-negative real number proportional to the gene’s transcriptional activity in that sample. After normalization and transformation, these values can be treated as coordinates in a real vector space.

Thus, each gene is represented as a vector $\vec{v} \in \mathbb{R}^n$, where n is the number of biological samples (e.g., tissues or time points), and each component $v_i \geq 0$ corresponds to the expression level in sample i . These vectors encode the distribution of gene activity across conditions and form the geometric basis for all downstream analyses.

Note on Stress Responses. In certain studies, a subset of the samples may correspond to stress conditions. Here, “stress” refers to any external or internal perturbation—such as mechanical damage, heat shock, pathogen infection, or nutrient deprivation—that triggers measurable changes in gene expression. These are often analyzed relative to baseline conditions using log fold-change transformations, such that the resulting vectors reflect differential expression patterns, including up- and downregulation across dimensions.

1.2 Normalization of Expression Values

Raw gene expression data is normalized in the following steps:

1. Division by gene-wise standard deviation for scale normalization.
2. Quantile normalization across samples to reduce global expression biases.
3. Log-transformation: $x \mapsto \log(1 + x)$ to stabilize variance and suppress outliers.
4. For stress response studies, log fold-changes are computed:

$$\Delta_{\text{stress}} = \log(1 + x_{\text{stress}}) - \log(1 + x_{\text{control}})$$

This effectively makes each stress related axis show the fold change in gene expression under stress, e.g. “drought in root divided by control root”.

1.3 Calculation of family Centroids

To summarize the expression profile of each family, Tensor Omics computes a centroid vector for every group. The centroid represents the average expression pattern of all genes assigned to the

family across all conditions or tissues.

There are two modes for selecting the set of genes to include in the centroid:

- **Full mode:** All genes assigned to the family are used to compute the centroid (mean vector).
- **Orthologs-only mode:** Only genes explicitly marked as orthologs are included in the centroid calculation. This mode is useful for focusing on conserved gene relationships across species.

1.4 Euclidean Distance Between Expression Vectors

The Euclidean distance is used throughout Tensor Omics to quantify the geometric difference between two gene expression vectors, or between a gene and its family centroid. Given two vectors $\vec{v}_1, \vec{v}_2 \in \mathbb{R}^n$, the Euclidean distance is defined as:

$$d(\vec{v}_1, \vec{v}_2) = \sqrt{\sum_{i=1}^n (v_{1,i} - v_{2,i})^2}$$

This metric measures the straight-line distance between two points in n -dimensional space. In the context of gene expression, it reflects the overall magnitude of difference in expression profiles across all samples or conditions.

For example, the distance between a paralog’s expression vector and the centroid of its gene family is used to detect outlier genes and to quantify the degree of divergence.

1.5 Outlier Detection Based on Family-Scaled Distances

To robustly identify genes with divergent expression profiles within their families, Tensor Omics implements an outlier detection procedure based on family-scaled Euclidean distances. The process consists of the following steps:

1. **Compute Euclidean Distances:** For each gene, calculate the Euclidean distance d_i between its expression vector and the centroid of its assigned family (see Section 1.4).
2. **Estimate Family-Specific Scaling Factors:** For each family f , we compute a family-specific scaling factor. This can be the maximum distance between any ortholog and the ortholog centroid. For families lacking sufficient orthologs, a LOESS smoothing strategy is used to estimate an appropriate scaling factor based on the family’s median intra-member distance.
3. **Relative Distance Index (RDI):** For each gene i in family f , compute the Relative Distance Index (RDI):

$$\text{RDI}_i = \frac{d_i}{d_{\text{scale}}}$$

This normalizes the gene’s distance by the expected intra-family variability, making the metric comparable across families of different sizes and variances.

4. **Outlier Detection:** Genes are flagged as outliers if their RDI exceeds a chosen percentile threshold (typically the 95th percentile) of the empirical RDI distribution. That is, gene i is an outlier if $\text{RDI}_i \geq T$, where T is the threshold value corresponding to the desired percentile.

This approach ensures that outlier detection is adaptive to the structure and variability of each gene family, and is robust to differences in family size and dispersion. The LOESS-based scaling is particularly important for families with few members.

1.6 Tissue Versatility and Specificity

Tissue Versatility quantifies how uniformly a gene is expressed across all tissues (or, more generally, axes of the expression space). It is defined as the cosine of the angle between a gene's expression vector $\vec{v}$ and the space diagonal $\vec{d} = (1, 1, \dots, 1)$. The smaller this angle (the larger the cosine), the greater the tissue (axis) versatility of the gene's expression.

Interpretation:

- A lower angle (cosine closer to 1) means the gene is more evenly expressed across all tissues (high versatility).
- A larger angle (cosine closer to 0) means the gene is active in only a few tissues (high specificity).

Dimensionality Dependence

The maximum possible angle to the diagonal depends on the number of tissues n . The cosine of this angle is minimal when a gene is expressed in only one tissue:

$$\cos(\varphi_{\max_angle}) = \frac{1}{\sqrt{n}}$$

Normalized Tissue Versatility

To make versatility comparable across datasets with different numbers of tissues, we define a normalized version that scales values into the interval $[0, 1]$:

$$\text{Tissue Versatility} = \frac{\cos(\varphi) - \frac{1}{\sqrt{n}}}{1 - \frac{1}{\sqrt{n}}}$$

This maps the cosine values into the interval $[0, 1]$:

- 1: perfectly uniform expression (maximum versatility)
- 0: expression in only one tissue (minimum versatility)

Tissue Specificity

Tissue specificity is defined as the complement of normalized tissue versatility:

$$\text{Tissue Specificity} = 1 - \text{Tissue Versatility}$$

Values close to 1 indicate high tissue-specificity; values close to 0 indicate broad, uniform expression.

2 Error Handling

In case you encounter errors while using TOX, the following error codes may help you identify the issue. The error codes are grouped into categories based on their nature, such as I/O errors, format validation errors, memory errors, and internal errors.

Table 1: TOX Error Codes Reference

Error Code	Description
0	Success
1xx: I/O / File Reading Errors	
101	Could not open file.

Continued on next page

Table 1 – continued from previous page

Error Code	Description
102	Could not read magic number.
103	Could not read array type code.
104	Could not read number of dimensions.
105	Could not read array dimensions.
106	Could not read character length.
107	Could not read array data.
112	Could not write magic number
113	Could not write array type code
114	Could not write number of dimensions
115	Could not write dimensions
116	Could not write character length
117	Could not write array data
2xx: Format / Input Validation Errors	
200	Invalid format detected (e.g., magic mismatch or corrupt format).
201	Invalid input provided.
202	Empty input arrays provided (e.g., zero-length arrays).
203	Dimension mismatch detected (shape inconsistencies).
204	NaN or Inf found in input data where not allowed.
205	Unsupported data type encountered (e.g., complex types).
206	Array size mismatch
208	Array index out of bounds
209	Division by zero detected
3xx: Memory Errors	
301	Memory allocation failed.
302	Null pointer reference encountered.
5xxx: Fortran Runtime / Unit State Errors	
5002	Fortran runtime error: unit not open or not connected.
9xxx: Internal / Unknown Errors	
9001	Internal error: unexpected state or invariant violated.
9999	Unknown error (fallback for unexpected errors).

3 Implementation Notes

This section details the practical pipeline for using the TOX library to read and process tabular data. The workflow is consistent across Fortran, Python, and R, starting with reading a CSV file into a staging area of strings and then converting columns to their proper data types.

3.1 Fortran Pipeline

3.1.1 Overview

1. Data dependencies and relationships All arrays read have dependencies to each other. The gene ids array stores the ids, column number seven of the expression vectors belongs to entry number seven in the gene ids. Entry number nine in the gene to family mapping array, belongs to entry nine in the gene ids. These dependencies must be given at any time, therefore, only the provided functions shall be used to modify arrays in any form. For error code details please check error handling section.
2. Error handling reference (placeholder for now)
3. General usage principle (Placeholder for now)

3.1.2 Data Management

1. Data Reading and Processing

(a) Reading Data

• Read Gene IDs from TSV File

Reads gene identifiers from a TSV file containing expression data. Typically reads from the first column after skipping header rows. Call the `read_gene_ids_from_tsv_file` subroutine with the following arguments:

Input arguments:

- `filename` character(len=*): Path to the TSV file
- `n_header_rows` integer(int32): Number of header rows to skip (typically 1)
- `gene_col` integer(int32): Column index containing gene IDs (typically 1)

Output arguments:

- `gene_ids` character(len=*): Pre-allocated array for gene identifiers
- `ierr` integer(int32): Error code

Listing 1: Read Gene IDs from TSV File

```
character(len=256) :: gene_ids(n_genes) ! Pre-allocate with
expected size
integer(int32) :: ierr

call read_gene_ids_from_tsv_file('expression_data.tsv', gene_ids,
1, 1, ierr)
```

• Read Expression Vectors

Reads expression data from one or more TSV/CSV files and populates a pre-allocated 2D array (samples x genes). Uses hashmap for efficient gene ID matching and validates for finite values. Call the `read_expression_vectors` subroutine with the following arguments:

Input arguments:

- `file_list` character(len=*): Array of file paths to read from
- `gene_ids` character(len=*): Reference gene IDs for validation and matching
- `expression_vectors` real(real64): Pre-allocated 2D array for expression data
- `n_header_rows` integer(int32): Number of header rows to skip
- `gene_col` integer(int32): Column index containing gene IDs
- `value_cols` integer(int32): Array of column indices containing expression values
- `start_row` integer(int32): Starting row index in expression_vectors array

– delimiter character(len=1), optional: Column delimiter (default: tab)

Output arguments:

- expression_vectors real(real64): Pre-allocated 2D array for expression data populated with data
- ierr integer(int32): Error code

Listing 2: Read Expression Vectors

```
! Pre-allocate arrays
real(real64) :: expr_data(n_samples, n_genes)

! Read TSV file with tab delimiter (default)
call read_expression_vectors(files, gene_ids, expr_data, 1, 1,
    value_cols, 1, ierr)

! Read CSV file with comma delimiter
call read_expression_vectors(files, gene_ids, expr_data, 1, 1,
    value_cols, 1, ierr, ',')
```

• **Read OrthoFinder Family Mapping**

Reads gene family assignments from OrthoFinder output TSV file and creates mapping arrays between genes and families. The `gene_to_fam` array contains for each gene the index of it's family. Uses hashmap for efficient gene lookup. Call the `read_orthofinder_file` subroutine with the following arguments:

Input arguments:

- filename character(len=*): Path to OrthoFinder TSV file
- gene_ids character(len=*): Pre-allocated array of gene IDs to match

Output arguments:

- family_ids character(len=*): Pre-allocated array for family identifiers
- gene_to_fam integer(int32): Pre-allocated array for gene to family mapping
- ierr integer(int32): Error code

Listing 3: Read OrthoFinder Family Mapping

```
character(len=256) :: family_ids(n_families) ! Pre-allocate with
    expected family count
integer(int32) :: gene_to_fam(n_genes) ! Pre-allocate with gene
    count
integer(int32) :: ierr

call read_orthofinder_file('Orthogroups.tsv', gene_ids, family_ids,
    gene_to_fam, ierr)
```

Important Implementation Details:

- All output arrays must be pre-allocated with the correct dimensions and sizes
- Uses hashmap data structure for O(1) gene ID lookups
- Validates for finite values (rejects NaN and Inf) in expression data
- Handles both TSV (tab-separated) and CSV (comma-separated) formats, `gene_to_fam` array uses 1-based indexing for family assignments (0 indicates unassigned)

Typical Data Reading Workflow:

Listing 4: Complete Data Reading Workflow

```
! 1. Pre-allocate arrays
```



```

! 2. Read gene IDs from expression file
call read_gene_ids_from_tsv_file('expression.tsv', gene_ids, 1, 1,
    ierr)

! 3. Read expression data (columns 2-68 contain sample data)
character(len=256) :: files(1) = ['expression.tsv']

call read_expression_vectors(files, gene_ids, expr_data, 1, 1,
    samples, 1, ierr)

! 4. Read family mapping from OrthoFinder
call read_orthofinder_file('Orthogroups.tsv', gene_ids, family_ids,
    gene_to_fam, ierr)

```

2. Data Filtering and Validation

• Get Unassigned Genes Mask

Creates a logical mask identifying genes that have no family assignment (gene_to_fam value = 0). Call the `get_unassigned_mask` subroutine with the following arguments:

Input arguments:

- gene_to_fam integer(int32): Gene to family mapping array

Output arguments:

- mask logical: Logical mask array (same size as gene_to_fam)
- n_genes_kept integer(int32): Number of genes with family assignments
- ierr integer(int32): Error code

Listing 5: Create Unassigned Genes Mask

```

logical :: unassigned_mask(n_genes)
integer(int32) :: n_kept, ierr

call get_unassigned_mask(gene_to_fam, unassigned_mask, n_kept)
! unassigned_mask(i) = .true. if gene i has family assignment
! n_kept = count of genes with family assignments

```

• Apply Unassigned Mask

Filters out genes without family assignments from all data arrays (gene_ids, expression_vectors, gene_to_fam). Arrays are reallocated to the filtered size. Call the `apply_unassigned_mask` subroutine with the following arguments:

Input arguments:

- mask logical: Logical mask from `get_unassigned_mask`
- n_genes_kept integer(int32): Number of genes to keep (from `get_unassigned_mask`)

Input/Output arguments:

- gene_ids character(len=*), allocatable: Gene IDs array (reallocated)
- expression_vectors real(real64), allocatable: Expression data matrix (reallocated)
- gene_to_fam integer(int32), allocatable: Gene to family mapping (reallocated)
- ierr integer(int32): Error code

Listing 6: Filter Unassigned Genes

```

character(len=256), allocatable :: gene_ids(n_genes)
real(real64), allocatable :: expr_data(n_samples, n_genes)
integer(int32), allocatable :: gene_to_fam(n_genes)

```

```

logical :: unassigned_mask(n_genes)
integer(int32) :: n_kept, ierr

! First create the mask
call get_unassigned_mask(gene_to_fam, unassigned_mask, n_kept)

! Then apply it to filter all arrays
call apply_unassigned_mask(gene_ids, expr_data, gene_to_fam, &
                           unassigned_mask, n_kept, ierr)
! Arrays are now reallocated to size n_kept

```

- **Validate Data Structure**

Performs comprehensive validation of data dimensions and consistency across all arrays. Call the `validate_data_structure` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Number of genes
- `n_families` integer(int32): Number of gene families
- `n_samples` integer(int32): Number of samples
- `gene_ids` character(len=*): Gene identifier array
- `family_ids` character(len=*): Family identifier array
- `gene_to_family` integer(int32): Gene to family mapping
- `expression` real(real64): Expression data matrix
- `centroids` real(real64): Family centroids matrix
- `shift_vectors` real(real64): Shift vectors matrix
- `ierr` integer(int32): Error code

Listing 7: Validate Data Structure

```

call validate_data_structure(n_genes, n_families, n_samples, gene_ids,
&
                           family_ids, gene_to_fam, expr_data, &
                           centroids, shift_vectors, ierr)

```

- **Validate Expression Data**

Validates expression data matrix for NaN, Inf, and optionally negative values. Call the `validate_expression_data` subroutine with the following arguments:

Input arguments:

- `expression` real(real64): Expression data matrix to validate
- `allow_negative` logical: Whether to allow negative expression values
- `ierr` integer(int32): Error code

Listing 8: Validate Expression Data

```

! Validate expression data (disallow negative values)
call validate_expression_data(expr_data, .false., ierr)

! Validate allowing negative values (e.g., for log-fold changes)
call validate_expression_data(expr_data, .true., ierr)

```

- **Validate Empty Strings**

Validates that a string array contains no empty or whitespace-only entries. Call the `validate_empty_strings` subroutine with the following arguments:

Input arguments:

- `string_array character(len=*)`: String array to validate
- `array_name character(len=*)`: Name for error messages

Output arguments:

- `ierr integer(int32)`: Error code

Listing 9: Validate String Array

```
call validate_empty_strings(gene_ids, "Gene IDs", ierr)
call validate_empty_strings(family_ids, "Family IDs", ierr)
```

- **Validate Gene to Family Mapping**

Validates that all gene to family indices are within valid range (1 to `n_families`). Call the `validate_gene_to_family_mapping` subroutine with the following arguments:

Input arguments:

- `gene_to_family integer(int32)`: Gene to family mapping array
- `n_families integer(int32)`: Number of valid families

Output arguments:

- `ierr integer(int32)`: Error code

Listing 10: Validate Gene-Family Mapping

```
call validate_gene_to_family_mapping(gene_to_fam, n_families, ierr)
```

- **Validate String Array Uniqueness**

Validates that all entries in a string array are unique (no duplicates). Uses a hashset implementation internally for $O(n)$ runtime. Call the `validate_string_array_uniqueness` subroutine with the following arguments:

Input arguments:

- `string_array character(len=*)`: String array to check for duplicates

Output arguments:

- `ierr integer(int32)`: Error code

Listing 11: Validate Unique Strings

```
call validate_string_array_uniqueness(gene_ids, ierr)
call validate_string_array_uniqueness(family_ids, ierr)
```

3. Data Accessors

- **Get Gene Index**

Finds the array index for a given gene identifier. Call the `get_gene_index` function with the following arguments:

Input arguments:

- `gene_ids character(len=*)`: Array of gene identifiers
- `gene_id character(len=*)`: Gene ID to find

Return value:

- `index integer(int32)`: Array index (1-based) or 0 if not found

Listing 12: Find Gene Index

```
integer(int32) :: gene_idx

gene_idx = get_gene_index(gene_ids, 'NP_001000001.1')
if (gene_idx > 0) then
    ! Gene found at position gene_idx
```

```

else
    ! Gene not found
end if

```

• Get Family Index

Finds the array index for a given family identifier. Call the `get_family_index` function with the following arguments:

Input arguments:

- `family_ids` character(len=*): Array of family identifiers
- `family_id` character(len=*): Family ID to find

Return value:

- `index` integer(int32): Array index (1-based) or 0 if not found

Listing 13: Find Family Index

```

integer(int32) :: fam_idx

fam_idx = get_family_index(family_ids, 'OG0000001')

```

• Get Family for Gene Index

Retrieves the family index for a given gene index. Call the `get_family_for_gene_index` subroutine with the following arguments:

Input arguments:

- `gene_index` integer(int32): Gene array index (1-based)
- `gene_to_family` integer(int32): Gene to family mapping array

Output arguments:

- `family_index` integer(int32): Family array index (1-based)
- `ierr` integer(int32): Error code

Listing 14: Get Family for Gene

```

integer(int32) :: gene_idx, fam_idx, ierr

gene_idx = get_gene_index(gene_ids, 'NP_001000001.1')
call get_family_for_gene_index(gene_idx, gene_to_fam, fam_idx, ierr)

```

• Get Family Centroid

Retrieves the centroid expression profile for a given family. Call the `get_family_centroid` subroutine with the following arguments:

Input arguments:

- `family_index` integer(int32): Family array index (1-based)
- `family_centroids` real(real64): Family centroids matrix

Output arguments:

- `ierr` integer(int32): Error code
- `centroid_vector` real(real64): Centroid expression profile

Listing 15: Get Family Centroid

```

real(real64) :: centroid_vector(n_samples)
integer(int32) :: fam_idx, ierr

fam_idx = get_family_index(family_ids, 'OG0000001')
call get_family_centroid(fam_idx, centroids, centroid_vector, ierr)

```

Complete Validation Workflow:

Listing 16: Complete Data Validation Workflow

```
integer(int32) :: ierr

! 1. Filter unassigned genes
logical :: unassigned_mask(n_genes)
integer(int32) :: n_kept
call get_unassigned_mask(gene_to_fam, unassigned_mask, n_kept)
call apply_unassigned_mask(gene_ids, expr_data, gene_to_fam, &
                           unassigned_mask, n_kept, ierr)

! 2. Validate individual components
call validate_empty_strings(gene_ids, "Gene IDs", ierr)
call validate_empty_strings(family_ids, "Family IDs", ierr)
call validate_string_array_uniqueness(gene_ids, ierr)
call validate_string_array_uniqueness(family_ids, ierr)
call validate_gene_to_family_mapping(gene_to_fam, n_families, ierr)
call validate_expression_data(expr_data, .false., ierr)

! 3. Comprehensive structure validation
call validate_data_structure(n_genes, n_families, n_samples, gene_ids,
                             &
                             family_ids, gene_to_fam, expr_data, &
                             centroids, shift_vectors, ierr)

! 4. Use accessors for data retrieval
integer(int32) :: gene_idx, fam_idx
real(real64) :: centroid(n_samples)

gene_idx = get_gene_index(gene_ids, 'NP_001000001.1')
call get_family_for_gene_index(gene_idx, gene_to_fam, fam_idx, ierr)
call get_family_centroid(fam_idx, centroids, centroid, ierr)
```

3.1.3 Data Persistence

1. Array Serialization and Deserialization Module

The array serialization and deserialization module provides comprehensive routines for efficiently storing and retrieving multi-dimensional arrays of various data types.

(a) General Architecture

The module supports five data types (integer, real, character, logical and complex) across dimensions 1 through 5. All routines follow a consistent naming convention and parameter structure for ease of use.

(b) Supported Functions

- `serialize_int_Nd(arr, filename, ierr)` - Serialize N-dimensional integer arrays.
- `deserialize_int_Nd(arr, filename, ierr)` - Deserialize N-dimensional integer arrays.
- `serialize_real_Nd(arr, filename, ierr)` - Serialize N-dimensional real arrays.
- `deserialize_real_Nd(arr, filename, ierr)` - Deserialize N-dimensional real arrays.
- `serialize_char_Nd(arr, filename, ierr)` - Serialize N-dimensional character arrays.
- `deserialize_char_Nd(arr, filename, ierr)` - Deserialize N-dimensional character arrays.

- `serialize_logical_Nd(arr, filename, ierr)` - Serialize N-dimensional logical arrays.
- `deserialize_logical_Nd(arr, filename, ierr)` - Deserialize N-dimensional logical arrays.
- `serialize_complex_Nd(arr, filename, ierr)` - Serialize N-dimensional complex arrays.
- `deserialize_complex_Nd(arr, filename, ierr)` - Deserialize N-dimensional complex arrays.
- `get_array_metadata(filename, dims, max_dims, ndims, ierr, [clen])` - Retrieve array metadata from files

Where N represents the array dimension (1-5).

(c) Detailed Parameter Specifications

Serialization Functions Parameters:

- **arr** (input): The array to be serialized.
Type: Allocatable array of type `integer(int32)`, `real(real64)`, `character(len=clen)`, (logical) or `complex`
Dimensions: 1D to 5D
Intent: in
Note: *clen* has to be the size of the arrays elements
- **filename** (input): The name of the binary file to create.
Type: `character(len=*)`
Intent: in
- **ierr** (output): Error code indicating operation status.
Type: `integer(int32)`
Intent: out
Values: 0 indicates success, non-zero values indicate specific error conditions

Listing 17: Example usage of a serialize function

```
call serialize_int_2d(array, "test_file.bin", ierr)
```

(d) Metadata Retrieval Function Parameters:

- **filename** (input): The name of the binary file to read metadata from.
Type: `character(len=*)`
Intent: in
- **dims** (output): Array containing the dimensions of the stored array.
Type: `integer(int32)`, `dimension(max_dims)`
Intent: out **Note:** Must be preallocated, recommended with size 5
- **max_dims** (input): Size of the *dims* array (typically 5 for maximum supported dimensions).
Type: `integer(int32)`
Intent: in
- **ndims** (output): Number of dimensions actually used by the array.
Type: `integer(int32)`
Intent: out
- **ierr** (output): Error code indicating operation status.
Type: `integer(int32)`
Intent: out
Values: 0 indicates success, non-zero values indicate specific error conditions
- **clen** (output, optional): For character arrays, the length of character elements.
Type: `integer(int32)`
Intent: out

Note: This parameter is only used and required for character arrays

Listing 18: Usage of metadata retrieval

```
call get_array_metadata(filename, dims, max_dims, ndims, ierr, clen)
```

Deserialization Functions Parameters:

- **arr** (output): The array to be populated with deserialized data.
Type: Allocatable array of type `integer(int32)`, `real(real64)`, `character(len=clen)`, `logical`, or `complex`
Dimensions: 1D to 5D
Intent: out
Note: The array must be allocated with correct dimensions and correct element length before calling deserialization routines
- **filename** (input): The name of the binary file to read.
Type: `character(len=*)`
Intent: in
- **ierr** (output): Error code indicating operation status.
Type: `integer(int32)`
Intent: out
Values: 0 indicates success, non-zero values indicate specific error conditions

Listing 19: Example usage of a deserialize function

```
call deserialize_int_2d(array, "test_file.bin", ierr)
```

Listing 20: Full example

```
integer(int32), allocatable :: kallisto_expression(:, :),  
    kallisto_expression_2(:, :)  
integer(int32) :: dims(5)  
integer(int32) :: ndims, ierr, max_dims  
character(len=64) :: fname  
  
!Assume array is filled with data  
  
call serialize_real_2d(kallisto_expression, "kallisto_data_v1.bin",  
    ierr)  
if(.not. is_ok(ierr)) error stop ierr  
  
!Reading the data again  
call get_array_metadata("kallisto_data_v1.bin", dims, max_dims, ndims,  
    ierr)  
if(.not. is_ok(ierr)) error stop ierr  
allocate(kallisto_expression_2(dims(1), dims(2)))  
  
call deserialize_real_2d(kallisto_expression_2, "kallisto_data_v1.bin"  
    , ierr)  
if(.not. is_ok(ierr)) error stop ierr
```

(e) Error Handling

All functions return an error code through the `ierr` parameter. The error handling utility `tox errors` provides functions to check and interpret these codes.

(f) Implementation Notes

- All routines are implemented in standard Fortran for interoperability
- The module uses allocatable arrays for flexible memory management
- File operations use Fortran's native unformatted I/O for efficiency
- Error handling is consistent across all functions for predictable behavior

2. Standard Archive Operations

• Save Tox Data Archive

Saves multiple data arrays to a ZIP archive with automatic manifest creation and temporary file management. Handles any combination of available standard data arrays. This function can only be used with standard gene expression data. If other data needs to be saved, see section "Non-standard archive operations". Call the `save_tox_data` subroutine with the following arguments:

Input arguments:

- `zip_filename` character(len=*): Output ZIP file name
- `ierr` integer(int32): Error code

Optional input arguments (provide both array and filename for each data type):

- `gene_ids` character(len=*): Gene IDs array
- `gene_ids_file` character(len=*): Internal filename for gene IDs
- `expression` real(real64): Expression data matrix
- `expression_file` character(len=*): Internal filename for expression data
- `gene_to_family` integer(int32): Gene to family mapping
- `gene_to_family_file` character(len=*): Internal filename for mapping
- `family_ids` character(len=*): Family IDs array
- `family_ids_file` character(len=*): Internal filename for family IDs
- `family_centroids` real(real64): Family centroids matrix
- `family_centroids_file` character(len=*): Internal filename for centroids
- `shift_vectors` real(real64): Shift vectors matrix
- `shift_vectors_file` character(len=*): Internal filename for shift vectors

Listing 21: Save Complete Data Archive

```
integer(int32) :: ierr

! Save all available data arrays
call save_tox_data("complete_data.zip", ierr, &
    gene_ids=gene_ids, gene_ids_file="genes.bin", &
    expression=expr_data, expression_file="expression.
        bin", &
    gene_to_family=gene_to_fam, gene_to_family_file="
        mapping.bin", &
    family_ids=family_ids, family_ids_file="families.bin
        ", &
    family_centroids=centroids, family_centroids_file="
        centroids.bin", &
    shift_vectors=shift_vecs, shift_vectors_file="shifts
        .bin")
```

Listing 22: Save Partial Data Archive

```
! Save only gene IDs and expression data
call save_tox_data("expression_only.zip", ierr, &
    gene_ids=gene_ids, gene_ids_file="genes.bin", &
```



```

        expression=expr_data, expression_file="expression.
        bin")

! Save only family-related data
call save_tox_data("families_only.zip", ierr, &
        family_ids=family_ids, family_ids_file="families.bin
        ", &
        family_centroids=centroids, family_centroids_file="
        centroids.bin")

```

• Read Tox Data Archive

Reads data arrays from a ZIP archive created with `save_tox_data`. Automatically handles array allocation based on archive contents. The function will try to read all requested arrays (optional arguments that are provided) from the given archive. If an array can not be found in the archive, the output arrays won't be allocated or returned. Call the `read_tox_data` subroutine with the following arguments:

Input arguments:

- `zip_filename` character(len=*): Input ZIP file name
- `ierr` integer(int32): Error code

Optional output arguments:

- `gene_ids` character(len=:), allocatable: Gene IDs array
- `expression` real(real64), allocatable: Expression data matrix
- `gene_to_family` integer(int32), allocatable: Gene to family mapping
- `family_ids` character(len=:), allocatable: Family IDs array
- `family_centroids` real(real64), allocatable: Family centroids matrix
- `shift_vectors` real(real64), allocatable: Shift vectors matrix

Optional filename output arguments:

- `gene_ids_file` character(len=:), allocatable: Internal filename used
- `expression_file` character(len=:), allocatable: Internal filename used
- `gene_to_family_file` character(len=:), allocatable: Internal filename used
- `family_ids_file` character(len=:), allocatable: Internal filename used
- `family_centroids_file` character(len=:), allocatable: Internal filename used
- `shift_vectors_file` character(len=:), allocatable: Internal filename used

Listing 23: Read Complete Data Archive

```

character(len=:), allocatable :: loaded_gene_ids(:), loaded_family_ids
(:)
real(real64), allocatable :: loaded_expr(:, :), loaded_centroids(:, :)
real(real64), allocatable :: loaded_shifts(:, :)
integer(int32), allocatable :: loaded_gene_to_fam(:)
integer(int32) :: ierr

! Read all available data from archive
call read_tox_data("complete_data.zip", ierr, &
        gene_ids=loaded_gene_ids, &
        expression=loaded_expr, &
        gene_to_family=loaded_gene_to_fam, &
        family_ids=loaded_family_ids, &
        family_centroids=loaded_centroids, &
        shift_vectors=loaded_shifts)

```

Listing 24: Read Partial Data Archive

```
! Read only specific arrays (others remain unallocated)
call read_tox_data("expression_only.zip", ierr, &
                  gene_ids=loaded_gene_ids, &
                  expression=loaded_expr)
```

Listing 25: Read with Filename Retrieval

```
character(len=:), allocatable :: actual_gene_file, actual_expr_file
integer(int32) :: ierr

! Get both data and internal filenames
call read_tox_data("complete_data.zip", ierr, &
                  gene_ids=loaded_gene_ids, gene_ids_file=
                    actual_gene_file, &
                  expression=loaded_expr, expression_file=
                    actual_expr_file)
```

3. Non-Standard Archive Operations

• Create Custom ZIP Archive

Creates a ZIP archive from arbitrary files with key-based manifest for non-standard data. Call the `create_zip_archive` subroutine with the following arguments:

Input arguments:

- `zip_filename` `character(len=*)`: Output ZIP file name
- `keys` `character(len=*)`: Array of keys for manifest entries
- `filenames` `character(len=*)`: Array of file paths to include
- `ierr` `integer(int32)`: Error code

Listing 26: Create Custom Archive

```
character(len=32) :: keys(3) = ['test_data_1', 'test_data_2', '
    test_data_3']
character(len=256) :: files(3) = ['test_data_1.bin', 'test_2.bin', '
    test_3.bin']
integer(int32) :: ierr

call create_zip_archive("custom_data.zip", keys, files, ierr)
```

• Extract ZIP Archive

Extracts all files from a ZIP archive to the disk in the current directory and returns key-value pairs from the manifest. Call the `extract_zip_archive` subroutine with the following arguments:

Input arguments:

- `zip_filename` `character(len=*)`: Input ZIP file name
- `ierr` `integer(int32)`: Error code

Output arguments:

- `keys` `character(len=:)`, `allocatable`: Array of keys from manifest
- `filenames` `character(len=:)`, `allocatable`: Array of extracted filenames

Listing 27: Extract Custom Archive

```
character(len=:), allocatable :: archive_keys(:), archive_files(:)
integer(int32) :: ierr
```

```
call extract_zip_archive("custom_data.zip", archive_keys,
    archive_files, ierr)
```

Complete Archive Workflow:

Listing 28: Complete Archive Management Workflow

```
integer(int32) :: ierr

! 1. Save processed data
call save_tox_data("analysis_results.zip", ierr, &
    gene_ids=gene_ids, gene_ids_file="genes.bin", &
    expression=expr_data, expression_file="expression.
        bin", &
    family_centroids=centroids, family_centroids_file="
        centroids.bin")

! 2. Later, read the data back
character(len=:), allocatable :: loaded_genes(:)
real(real64), allocatable :: loaded_expr(:, :), loaded_centroids(:, :)
call read_tox_data("analysis_results.zip", ierr, &
    gene_ids=loaded_genes, &
    expression=loaded_expr, &
    family_centroids=loaded_centroids)

! 3. Create custom archive with additional results
character(len=32) :: custom_keys(2) = ['data_1', 'data_2']
character(len=256) :: custom_files(2) = ['data_1_v2.bin', 'data_2_v1.
    bin']
call create_zip_archive("full_analysis.zip", custom_keys, custom_files
    , ierr)

! 4. Extract and verify custom archive
character(len=:), allocatable :: extracted_keys(:), extracted_files(:)
call extract_zip_archive("full_analysis.zip", extracted_keys,
    extracted_files, ierr)
```

Key Features of Archive Operations:

- **Automatic Memory Management:** Arrays automatically allocated to correct sizes
- **Flexible Data Selection:** Save/read any combination of data arrays
- **Cross-Platform Compatibility:** Archives work across R, Python, and Fortran implementations
- **Manifest-Based:** Automatic tracking of file contents and relationships
- **Temporary File Handling:** Automatic cleanup of intermediate files
- **Error Recovery:** Comprehensive error checking and reporting

Archive Structure:

```
archive.zip
|- manifest.txt          (key=filename mappings)
|- genes.bin             (gene IDs, if provided)
|- expression.bin        (expression data, if provided)
|- mapping.bin           (gene to family, if provided)
|- families.bin          (family IDs, if provided)
```

```

|- centroids.bin          (family centroids, if provided)
|- shifts.bin             (shift vectors, if provided)

```

4. Individual Array Save/Load Operations

• Save Gene IDs

Saves a gene IDs array to a binary file with automatic character length preservation. Call the `save_gene_ids` subroutine with the following arguments:

Input arguments:

- `gene_ids` `character(len=*)`: Array of gene identifiers
- `filename` `character(len=*)`: Output binary file name
- `ierr` `integer(int32)`: Error code

Listing 29: Save Gene IDs Array

```

character(len=256) :: gene_ids(10000)
integer(int32) :: ierr

call save_gene_ids(gene_ids, 'gene_ids.bin', ierr)

```

• Load Gene IDs

Loads a gene IDs array from a binary file with automatic allocation and character length detection. Call the `load_gene_ids` subroutine with the following arguments:

Input arguments:

- `filename` `character(len=*)`: Input binary file name
- `ierr` `integer(int32)`: Error code

Output arguments:

- `gene_ids` `character(len=:)`, `allocatable`: Loaded gene IDs array

Listing 30: Load Gene IDs Array

```

character(len=:), allocatable :: loaded_gene_ids(:)
integer(int32) :: ierr

call load_gene_ids(loaded_gene_ids, 'gene_ids.bin', ierr)
! loaded_gene_ids automatically allocated with correct size and
  character length

```

• Save Expression Vectors

Saves an expression data matrix to a binary file. Call the `save_expression_vectors` subroutine with the following arguments:

Input arguments:

- `expression` `real(real64)`: 2D expression data matrix
- `filename` `character(len=*)`: Output binary file name
- `ierr` `integer(int32)`: Error code

Listing 31: Save Expression Data

```

real(real64) :: expr_data(n_samples, n_genes)
integer(int32) :: ierr

call save_expression_vectors(expr_data, 'expression.bin', ierr)

```

• Load Expression Vectors

Loads an expression data matrix from a binary file with automatic allocation. Call the `load_expression_vectors` subroutine with the following arguments:

Input arguments:

- filename character(len=*): Input binary file name
- ierr integer(int32): Error code

Output arguments:

- expression real(real64), allocatable: Loaded expression data matrix

Listing 32: Load Expression Data

```
real(real64), allocatable :: loaded_expr(:, :)
integer(int32) :: ierr

call load_expression_vectors(loaded_expr, 'expression.bin', ierr)
! loaded_expr automatically allocated as (samples x genes)
```

- **Save Gene to Family Mapping**

Saves a gene to family mapping array to a binary file. Call the `save_gene_to_family` subroutine with the following arguments:

Input arguments:

- gene_to_family integer(int32): Gene to family mapping array
- filename character(len=*): Output binary file name
- ierr integer(int32): Error code

Listing 33: Save Gene-Family Mapping

```
integer(int32) :: gene_to_fam(n_genes)
integer(int32) :: ierr

call save_gene_to_family(gene_to_fam, 'gene_to_fam.bin', ierr)
```

- **Load Gene to Family Mapping**

Loads a gene to family mapping array from a binary file with automatic allocation. Call the `load_gene_to_family` subroutine with the following arguments:

Input arguments:

- filename character(len=*): Input binary file name
- ierr integer(int32): Error code

Output arguments:

- gene_to_family integer(int32), allocatable: Loaded gene to family mapping

Listing 34: Load Gene-Family Mapping

```
integer(int32), allocatable :: loaded_gene_to_fam(:)
integer(int32) :: ierr

call load_gene_to_family(loaded_gene_to_fam, 'gene_to_fam.bin', ierr)
```

- **Save Family IDs**

Saves a family IDs array to a binary file with automatic character length preservation. Call the `save_family_ids` subroutine with the following arguments:

Input arguments:

- family_ids character(len=*): Array of family identifiers
- filename character(len=*): Output binary file name
- ierr integer(int32): Error code

Listing 35: Save Family IDs

```

character(len=256) :: family_ids(n_families)
integer(int32) :: ierr

call save_family_ids(family_ids, 'family_ids.bin', ierr)

```

- **Load Family IDs**

Loads a family IDs array from a binary file with automatic allocation and character length detection. Call the `load_family_ids` subroutine with the following arguments:

Input arguments:

- filename character(len=*): Input binary file name
- ierr integer(int32): Error code

Output arguments:

- family_ids character(len=:), allocatable: Loaded family IDs array

Listing 36: Load Family IDs

```

character(len=:), allocatable :: loaded_family_ids(:)
integer(int32) :: ierr

call load_family_ids(loaded_family_ids, 'family_ids.bin', ierr)

```

- **Save Family Centroids**

Saves a family centroids matrix to a binary file. Call the `save_family_centroids` subroutine with the following arguments:

Input arguments:

- family_centroids real(real64): 2D family centroids matrix
- filename character(len=*): Output binary file name
- ierr integer(int32): Error code

Listing 37: Save Family Centroids

```

real(real64) :: centroids(67, 15512)
integer(int32) :: ierr

call save_family_centroids(centroids, 'centroids.bin', ierr)

```

- **Load Family Centroids**

Loads a family centroids matrix from a binary file with automatic allocation. Call the `load_family_centroids` subroutine with the following arguments:

Input arguments:

- filename character(len=*): Input binary file name
- ierr integer(int32): Error code

Output arguments:

- family_centroids real(real64), allocatable: Loaded family centroids matrix

Listing 38: Load Family Centroids

```

real(real64), allocatable :: loaded_centroids(:, :)
integer(int32) :: ierr

call load_family_centroids(loaded_centroids, 'centroids.bin', ierr)

```

- **Save Shift Vectors**

Saves a shift vectors matrix to a binary file. Call the `save_shift_vectors` subroutine with the following arguments:

Input arguments:

- `shift_vectors` `real(real64)`: 2D shift vectors matrix
- `filename` `character(len=*)`: Output binary file name
- `ierr` `integer(int32)`: Error code

Listing 39: Save Shift Vectors

```
real(real64) :: shift_vecs(134, 88327) ! 2xsamples x genes
integer(int32) :: ierr

call save_shift_vectors(shift_vecs, 'shift_vectors.bin', ierr)
```

- **Load Shift Vectors**

Loads a shift vectors matrix from a binary file with automatic allocation. Call the `load_shift_vectors` subroutine with the following arguments:

Input arguments:

- `filename` `character(len=*)`: Input binary file name
- `ierr` `integer(int32)`: Error code

Output arguments:

- `shift_vectors` `real(real64)`, `allocatable`: Loaded shift vectors matrix

Listing 40: Load Shift Vectors

```
real(real64), allocatable :: loaded_shifts(:, :)
integer(int32) :: ierr

call load_shift_vectors(loaded_shifts, 'shift_vectors.bin', ierr)
```

- **Get Array Metadata**

Retrieves metadata (dimensions, character length) from a binary file without loading the entire array. Call the `get_array_metadata` subroutine with the following arguments:

Input arguments:

- `filename` `character(len=*)`: Binary file name to inspect
- `max_dims` `integer(int32)`: Maximum number of dimensions to return
- `ierr` `integer(int32)`: Error code

Output arguments:

- `dims` `integer(int32)`: Array dimensions
- `ndims` `integer(int32)`: Number of dimensions
- `char_len` `integer(int32)`, optional: Character length for string arrays

Listing 41: Inspect Array File Metadata

```
integer(int32) :: dims(5), ndims, char_len, ierr

! Get metadata for gene IDs file
call get_array_metadata('gene_ids.bin', dims, 5, ndims, ierr, char_len)
if (is_ok(ierr)) then
    print *, "Dimensions:", dims(1:ndims)
    print *, "Character length:", char_len
end if

! Get metadata for expression data file
call get_array_metadata('expression.bin', dims, 5, ndims, ierr)
```

```

if (is_ok(ierr)) then
    print *, "Expression data shape:", dims(1), "x", dims(2)
end if

```

Complete Individual Array Workflow:

Listing 42: Individual Array Save/Load Workflow

```

integer(int32) :: ierr

! 1. Save individual arrays
call save_gene_ids(gene_ids, 'gene_ids.bin', ierr)
call save_expression_vectors(expr_data, 'expression.bin', ierr)
call save_gene_to_family(gene_to_fam, 'mapping.bin', ierr)
call save_family_ids(family_ids, 'families.bin', ierr)
call save_family_centroids(centroids, 'centroids.bin', ierr)
call save_shift_vectors(shift_vecs, 'shifts.bin', ierr)

! 2. Check file metadata before loading
integer(int32) :: dims(5), ndims, char_len
call get_array_metadata('expression.bin', dims, 5, ndims, ierr)
if (is_ok(ierr) .and. ndims == 2) then
    print *, "Will load expression data:", dims(1), "samples x", dims
        (2), "genes"
end if

! 3. Load arrays (automatic allocation)
character(len=:), allocatable :: loaded_genes(:)
real(real64), allocatable :: loaded_expr(:, :)
integer(int32), allocatable :: loaded_mapping(:)
character(len=:), allocatable :: loaded_families(:)
real(real64), allocatable :: loaded_centroids(:, :)
real(real64), allocatable :: loaded_shifts(:, :)

call load_gene_ids(loaded_genes, 'gene_ids.bin', ierr)
call load_expression_vectors(loaded_expr, 'expression.bin', ierr)
call load_gene_to_family(loaded_mapping, 'mapping.bin', ierr)
call load_family_ids(loaded_families, 'families.bin', ierr)
call load_family_centroids(loaded_centroids, 'centroids.bin', ierr)
call load_shift_vectors(loaded_shifts, 'shifts.bin', ierr)

! 4. Verify loaded data
print *, "Loaded", size(loaded_genes), "genes"
print *, "Loaded expression data:", size(loaded_expr, 1), "x", size(
    loaded_expr, 2)
print *, "Loaded", size(loaded_families), "families"

```

Manual Archive Creation with Individual Arrays:

Listing 43: Manual Archive Creation

```

integer(int32) :: ierr

! 1. Save arrays to individual files
call save_gene_ids(gene_ids, 'temp_genes.bin', ierr)
call save_expression_vectors(expr_data, 'temp_expr.bin', ierr)

```



```

! 2. Create custom archive with specific keys
character(len=32) :: keys(2) = ['gene_data', 'expression_data']
character(len=256) :: files(2) = ['temp_genes.bin', 'temp_expr.bin']
call create_zip_archive('manual_archive.zip', keys, files, ierr)

! 3. Clean up temporary files (optional)
call delete_file('temp_genes.bin', ierr)
call delete_file('temp_expr.bin', ierr)

```

Key Features of Individual Array Operations:

- **Automatic Allocation:** Output arrays automatically allocated to correct sizes
- **Character Length Preservation:** String arrays maintain original character lengths
- **Metadata Inspection:** Check array properties without loading data
- **Binary Format:** Efficient storage for large datasets
- **Error Handling:** Comprehensive error checking for file operations
- **Memory Efficient:** Can process arrays larger than available memory

Performance Considerations:

- Binary format provides fast I/O for large arrays
- Automatic dimension detection eliminates manual size tracking
- Character length preservation avoids fixed-width string limitations
- Suitable for datasets with millions of genes and thousands of samples

3.1.4 Data Processing

1. Normalization of Expression Values

The subroutine runs a complete normalization pipeline on gene expression data. It first applies standard deviation normalization, then quantile normalization, followed by replicate averaging, and finally a log2 transformation. The result is returned in the `buf_log` array, which contains the fully normalized expression values. Additionally intermediate results of the performed calculations are available in output buffer arrays. Call the `normalization_pipeline` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Total number of genes
- `n_tissues` integer(int32): Number of axes (tissues / dimensions)
- `input_matrix` real(real64): Expression vector matrix with size `n_genes × n_tissues`
- `max_stack` integer(int32): Stack size for quicksort
- `group_s` integer(int32): Start column index for each replicate group (`n_grps`)
- `group_c` integer(int32): Number of columns per replicate group (`n_grps`)
- `n_grps` integer(int32): Number of replicate groups

Output arguments:

- `buf_stddev` real(real64): Buffer for std dev normalization (`n_genes * n_tissues`)
- `buf_quant` real(real64): Buffer for quantile normalization (`n_genes * n_tissues`)
- `buf_avg` real(real64): Buffer for replicate averaging (`n_genes * n_grps`)
- `buf_log` real(real64): Buffer for log2(x+1) transformation (`n_genes * n_grps`)
- `temp_col` real(real64): Temporary column vector for sorting (`n_genes`)
- `rank_means` real(real64): Buffer for rank means (`n_genes`)
- `perm` integer(int32): Permutation vector for sorting (`n_genes`)
- `stack_left` integer(int32): Left stack for quicksort (`max_stack`)

- `stack_right` integer(int32): Right stack for quicksort (`max_stack`)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 44: Normalization Pipeline

```
call normalization_pipeline(n_genes, n_tissues, input_matrix, buf_stddev
, buf_quant, buf_avg, buf_log, temp_col, rank_means, perm, stack_left
, stack_right, max_stack, group_s, group_c, n_grps, ierr)
```

Alternatively you can calculate each step of the normalization pipeline manually by calling the following subroutines in sequence:

(a) **Normalize by Standard Deviation**

The subroutine normalizes each gene's expression vector using `sqrt(mean(x^2))` across tissues (not classical standard deviation). Call the `normalize_by_std_dev` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Number of genes (rows)
- `n_tissues` integer(int32): Number of tissues (columns)
- `input_matrix` real(real64): Flattened input matrix of gene expression values (column-major)

Output arguments:

- `output_matrix` real(real64): Output normalized matrix (same shape as input)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 45: Standard Deviation Normalization

```
call normalize_by_std_dev(n_genes, n_tissues, input_matrix,
output_matrix, ierr)
```

(b) **Quantile Normalization**

The subroutine performs quantile normalization of a gene expression matrix (F42-compliant) and computes average expression per rank across tissues. Call the `quantile_normalization` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Number of genes (rows)
- `n_tissues` integer(int32): Number of tissues (columns)
- `input_matrix` real(real64): Flattened input matrix (column-major)
- `max_stack` integer(int32): Stack size passed from R

Output arguments:

- `output_matrix` real(real64): Output normalized matrix (same shape as input)
- `temp_col` real(real64): Temporary vector for column sorting (size `n_genes`)
- `rank_means` real(real64): Preallocated vector to store rank means (size `n_genes`)
- `perm` integer(int32): Permutation vector (size `n_genes`)
- `stack_left` integer(int32): Manual quicksort stack ($\geq \log_2(n_genes) + 10$)
- `stack_right` integer(int32): Manual quicksort stack (same size as `stack_left`)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 46: Quantile Normalization

```
call quantile_normalization(n_genes, n_tissues, input_matrix,
output_matrix, temp_col, rank_means, perm, stack_left, stack_right
, max_stack, ierr)
```

(c) **Calculate Tissue Averages**

The subroutine calculates tissue averages by averaging replicates within each group. For each group of tissue replicates, it computes the average expression per gene. The input matrix is column-major, flattened as a 1D array. Call the `calc_tiss_avg` subroutine with the following arguments:

Input arguments:

- `n_gene` integer(int32): Number of genes (rows)
- `n_grps` integer(int32): Number of tissue groups
- `group_s` integer(int32): Start column index for each group (length: `n_grps`)
- `group_c` integer(int32): Number of columns per group (length: `n_grps`)
- `input_matrix` real(real64): Flattened input matrix (length: `n_gene * sum(group_c)`)

Output arguments:

- `output_matrix` real(real64): Flattened output matrix (length: `n_gene * n_grps`)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 47: Tissue Averages Calculation

```
call calc_tiss_avg(n_gene, n_grps, group_s, group_c, input_matrix,
                  output_matrix, ierr)
```

(d) **Log2 Transformation**

The subroutine applies $\log_2(x + 1)$ transformation to each element of the input matrix. This is computed via $\log(x + 1) / \log(2)$ for compatibility with WebAssembly (WASM). Call the `log2_transformation` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Number of genes (rows)
- `n_tissues` integer(int32): Number of tissues (columns)
- `input_matrix` real(real64): Flattened input matrix (size: `n_genes * n_tissues`)

Output arguments:

- `output_matrix` real(real64): Output matrix (same size as input)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 48: Log2 Transformation

```
call log2_transformation(n_genes, n_tissues, input_matrix,
                        output_matrix, ierr)
```

2. Calculate Fold Change

If fold change calculation is needed, this subroutine can be used after normalization. The subroutine calculates $\log_2$ fold changes between condition and control columns. For each control-condition pair, this subroutine computes the $\log_2$ fold change by subtracting the expression value in the control column from the corresponding value in the condition column, for all genes. Call the `calc_fchange` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Total number of genes
- `n_cols` integer(int32): Number of columns in the input matrix
- `n_pairs` integer(int32): Number of control-condition pairs
- `control_cols` integer(int32): Control column indices (length `n_pairs`)
- `cond_cols` integer(int32): Condition column indices (length `n_pairs`)
- `i_matrix` real(real64): Input matrix, flattened (length: `n_genes × n_cols`)

Output arguments:

- `o_matrix` `real(real64)`: Output matrix for fold changes (length: `n_genes` × `n_pairs`)
- `ierr` `integer(int32)`: Error code (for error details please check error handling section)

Listing 49: Fold Change Calculation

```
call calc_fchange(n_genes, n_cols, n_pairs, control_cols, cond_cols,
  i_matrix, o_matrix, ierr)
```

3.1.5 Family-Based Analysis

(a) Calculation of family Centroids

The subroutine computes the centroid (mean expression vector) for each gene family in a gene expression dataset. It iterates over all families, selects the relevant genes for each family and calculates the mean vector for each family. There are two possible modes to use. To use all genes for calculation, select `all` mode. If you want to specify relevant genes for calculation, select `orthologs` mode. It is important that you provide a logical `ortholog_set` array when using `orthologs` mode. Call the `group_centroid` subroutine with the following arguments:

Input arguments:

- `expression_vectors` `real(real64)`: Expression vector matrix with size `n_axes` × `n_genes`
- `n_axes` `integer(int32)`: Number of axes (tissues / dimensions)
- `n_genes` `integer(int32)`: Total number of genes
- `gene_to_family` `integer(int32)`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)
- `n_families` `integer(int32)`: Total number of gene families
- `mode` `character(len=*)`: Calculation mode ("all" or "orthologs")
- (Optional) `ortholog_set` `logical`: Logical array indicating if a gene is part of the calculation subset (only used in `orthologs` mode)

Output arguments:

- `centroid_matrix` `real(real64)`: Matrix of calculated family centroid vectors with size `n_axes` × `n_families`
- `selected_indices` `integer(int32)`: Array of gene indices used for the calculation with length of `n_genes`
- `ierr` `integer(int32)`: Error code (for error details please check error handling section)

Listing 50: Calculation of Family Centroids

```
call group_centroid(expression_vectors, n_axes, n_genes,
  gene_to_family, n_families, centroid_matrix, mode,
  selected_indices, ierr, ortholog_set)
```

(b) Calculate shift vectors

Placeholder for now

(c) Calculate Mean Vector

Alternatively you can calculate a single mean vector for a given set of vectors by calling the `mean_vector` subroutine directly with the following arguments:

Input arguments:

- `expression_vectors` `real(real64)`: Expression vector matrix with size `n_axes` × `n_genes`
- `n_axes` `integer(int32)`: Number of axes (tissues / dimensions)
- `n_genes` `integer(int32)`: Total number of genes
- `gene_indices` `integer(int32)`: An array containing the column indices of the selected genes in `expression_vectors`

- `n_selected_genes integer(int32)`: Total number of genes in the current family to be calculated

Output arguments:

- `centroid real(real64)`: Calculated mean vector (centroid) with length of `n_axes`
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 51: Mean Vector Calculation

```
call mean_vector(expression_vectors, n_axes, n_genes, gene_indices,
                 n_selected_genes, centroid, ierr)
```

(d) **Calculate Distance to Centroid**

For each gene in the expression vector matrix, the subroutine computes the Euclidean distance to its corresponding family centroid and returns a distances array containing the distance value for each expression vector to its family centroid. Call the `distance_to_centroid` subroutine with the following arguments:

Input arguments:

- `n_genes integer(int32)`: Total number of genes
- `n_families integer(int32)`: Total number of gene families
- `genes real(real64)`: Gene expression matrix ($d \times n_genes$)
- `centroids real(real64)`: Family centroid matrix ($d \times n_genes$)
- `gene_to_fam integer(int32)`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)
- `d integer(int32)`: Expression vector dimension

Output arguments:

- `distances real(real64)`: Array of distances for each gene with length of `n_genes`

Listing 52: Distance to Centroid Calculation

```
call distance_to_centroid(n_genes, n_families, genes, centroids,
                         gene_to_fam, distances, d)
```

i. **Calculate Euclidean Distance**

Alternatively you can calculate a single Euclidean distance between two vectors manually by calling `euclidean_distance` subroutine with the following arguments:

Input arguments:

- `vec1 real(real64)`: First expression vector as array with length of dimension `d`
- `vec2 real(real64)`: Second expression vector as array with length of dimension `d`
- `d integer(int32)`: Dimension value

Output arguments:

- `result real(real64)`: Euclidean distance value between `vec1` and `vec2`

Listing 53: Euclidean distance calculation

```
call euclidean_distance(vec1, vec2, d, result)
```

(e) **Detect Gene Outliers**

The subroutine first computes family scaling, then calculates relative distance indices (RDI) and identifies outliers based on the specified percentile threshold. It returns a boolean array indicating whether each gene is an outlier. Call the `detect_outliers` subroutine with the following arguments:

Input arguments:

- `n_genes integer(int32)`: Total number of genes

- `n_families` integer(int32): Total number of gene families
- `distances` real(real64): Array of distances for each gene to its centroid with length of `n_genes`
- `gene_to_fam` integer(int32): Gene-to-family mapping array with the length of `n_genes` (1-based indexing)
- **(Optional)** `percentile` real(real64): Percentile threshold value for outlier detection (default is 95)

Input / Output arguments:

- `work_array` real(real64): Work array for sorting with length of `n_genes`
- `perm` integer(int32): Permutation array for sorting with length of `n_genes`
- `stack_left` integer(int32): Stack array for left indices during sorting with length of `n_genes`
- `stack_right` integer(int32): Stack array for right indices during sorting with length of `n_genes`
- `loess_x` real(real64): Reference x-coordinates for loess smoothing with length of `n_families`
- `loess_y` real(real64): Reference y-coordinates for loess smoothing with length of `n_families`
- `loess_n` integer(int32): Indices of reference points used for smoothing with length of `n_families`

Output arguments:

- `is_outlier` logical: Boolean array with length of `n_genes` indicating whether each gene is an outlier
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 54: Gene Outlier Calculation

```
call detect_outliers(n_genes, n_families, distances, gene_to_fam,
    work_array, perm, stack_left, stack_right, is_outlier, loess_x,
    loess_y, loess_n, ierr, percentile)
```

Alternatively you can calculate each step of `detect_outliers` manually with the following subroutine:

i. **Compute Family Scaling**

The subroutine calculates a scaling factor for each gene family by computing the median and standard deviation of gene-to-centroid distances within each family, then applies LOESS smoothing to these values. Call the `compute_family_scaling` subroutine with the following arguments:

Input arguments:

- `n_genes` integer(int32): Total number of genes
- `n_families` integer(int32): Total number of gene families
- `distances` real(real64): Array of distances for each gene to its centroid with length of `n_genes`
- `gene_to_fam` integer(int32): Gene-to-family mapping array with the length of `n_genes` (1-based indexing)

Input / Output arguments:

- `loess_x` real(real64): Reference x-coordinates for loess smoothing with length of `n_families`
- `loess_y` real(real64): Reference y-coordinates for loess smoothing with length of `n_families`
- `indices_used` integer(int32): Indices of reference points used for smoothing with length of `n_families`

- `perm_tmp integer(int32)`: Permutation array for sorting with length of `n_genes`

Output arguments:

- `dscale real(real64)`: Array of scaling factors per family with length of `n_families`
- `family_distances real(real64)`: Temporary work array with length of `n_genes` to store distances of genes within each family during calculation
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 55: Compute Family Scaling

```
call compute_family_scaling(n_genes, n_families, distances,
    gene_to_fam, dscale, loess_x, loess_y, indices_used, perm_tmp,
    stack_left_tmp, stack_right_tmp, family_distances, ierr)
```

ii. **Compute Family Scaling Helper**

This helper subroutine allocates internal arrays for easier usage and calls `compute_family_scaling`. Call the `compute_family_scaling_alloc` subroutine with the following arguments:

Input arguments:

- `n_genes integer(int32)`: Total number of genes
- `n_families integer(int32)`: Total number of gene families
- `distances real(real64)`: Array of distances for each gene to its centroid with length of `n_genes`
- `gene_to_fam integer(int32)`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)

Input / Output arguments:

- `loess_x real(real64)`: Reference x-coordinates for loess smoothing with length of `n_families`
- `loess_y real(real64)`: Reference y-coordinates for loess smoothing with length of `n_families`
- `indices_used integer(int32)`: Indices of reference points used for smoothing with length of `n_families`

Output arguments:

- `dscale real(real64)`: Array of scaling factors per family with length of `n_families`
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 56: Compute Family Scaling

```
call compute_family_scaling_alloc(n_genes, n_families, distances,
    gene_to_fam, dscale, loess_x, loess_y, indices_used, ierr)
```

iii. **Compute Relative Distance Index (RDI)**

The subroutine calculates the Relative Distance Index (RDI) for each gene by dividing its Euclidean distance to the family centroid by the scaling factor of its family. Call the `compute_rdi` subroutine with the following arguments:

Input arguments:

- `n_genes integer(int32)`: Total number of genes
- `distances real(real64)`: Array of distances for each gene to its centroid with length of `n_genes`
- `gene_to_fam integer(int32)`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)

- `dscale real(real64)`: Array of scaling factors per family with length of `n_families`

Input / Output arguments:

- `sorted_rdi real(real64)`: Work array for sorting with length of `n_genes`
- `perm integer(int32)`: Pre-initialized (1: `n_genes`) permutation array for sorting with length of `n_genes`
- `stack_left integer(int32)`: Stack array for sorting with length of `n_genes`
- `stack_right integer(int32)`: Stack array for sorting with length of `n_genes`

Output arguments:

- `rdi real(real64)`: Array of RDI values for each gene with length of `n_genes`

Listing 57: Compute Relative Distance Index (RDI)

```
call compute_rdi(n_genes, distances, gene_to_fam, dscale, rdi,
  sorted_rdi, perm, stack_left, stack_right)
```

iv. **Identify Outliers**

The subroutine identifies outliers based on the top percentile of RDI values. Call the `identify_outliers` subroutine with the following arguments:

Input arguments:

- `n_genes integer(int32)`: Total number of genes
- `rdi real(real64)`: Array of RDI values for each gene with length of `n_genes`
- `sorted_rdi real(real64)`: Array of sorted RDI values for each gene with length of `n_genes` - **this array must be filtered to remove negatives and be sorted in ascending order!**
- **(Optional)** `percentile real(real64)`: Percentile threshold value for outlier detection (default is 95)

Output arguments:

- `is_outlier logical`: Boolean array with length of `n_genes` indicating whether each gene is an outlier
- `threshold`: Actual threshold value used for outlier detection

Listing 58: Identify Outliers

```
call identify_outliers(n_genes, rdi, sorted_rdi, is_outlier,
  threshold, percentile)
```

(f) **Compute Normalized Tissue Versatility**

The subroutine calculates a normalized tissue versatility score for selected gene expression vectors. It is based on the angle between each gene expression vector and the space diagonal. Versatility is normalized to [0, 1], where 0 means uniform expression 1 means expression in only one axis. Call the `compute_tissue_versatility` subroutine with the following parameters:

Input arguments:

- `n_genes integer(int32)`: Total number of genes
- `n_vectors integer(int32)`: Total number of expression vectors
- `n_selected_axes integer(int32)`: Total number of selected axes
- `n_selected_vectors integer(int32)`: Total number of selected vectors
- `expression_vectors real(real64)`: Expression vector matrix with size `n_axes × n_vectors`
- `exp_vecs_selection_index logical`: Logical array with length of `n_vectors` to specify which vectors should be included in the calculation
- `axes_selection logical`: Logical array with length of `n_axes` to specify which axes should be included in the calculation

Output arguments:

- `tissue_versatilities` `real(real64)`: Array with the length of `n_selected_vectors` containing the calculated tissue versatilities
- `tissue_angles_deg` `real(real64)`: Array with the length of `n_selected_vectors` containing the calculated angles in degrees
- `ierr` `integer(int32)`: Error code (for error details please check error handling section)

Listing 59: Tissue Versatility Calculation

```
call compute_tissue_versatility(n_axes, n_vectors, expression_vectors,
    exp_vecs_selection_index, n_selected_vectors, axes_selection,
    n_selected_axes, tissue_versatilities, tissue_angles_deg, ierr)
```

3.1.6 Relative Axis Plane Projection

Placeholder

3.1.7 Contribution Analysis**1. Fortran Core Contribution Analysis Functions**

The Fortran module provides core routines for trajectory contribution analysis using cosine similarity and spike detection.

(a) Trajectory Contribution Calculation

Calculates integrated trajectory contribution using cosine similarity between factor and dependent vectors.

Mathematical Formula:

$$\text{contribution} = \frac{\text{factor} \cdot \text{dependent}}{\|\text{factor}\| \cdot \|\text{dependent}\|}$$

Arguments:

- `n_timepoints` `integer(int32)`: Number of timepoints
- `factor` `real(real64)`: Vector of independent variable [`n_timepoints`]
- `dependent` `real(real64)`: Vector of dependent variable [`n_timepoints`]
- `mode` `integer(int32)`: Processing mode (1 = Normal, 2 = RAP)
- `contribution` `real(real64)`: Output contribution value
- `ierr` `integer(int32)`: Error code

Listing 60: Trajectory Contribution Calculation

```
call trajectory_contribution(factor, dependent, n_timepoints, mode,
    contribution, ierr)
```

(b) Spike Contribution Calculation

Calculates timepoint-wise spike contributions using cosine similarity for each individual timepoint.

Mathematical Formula:

$$\text{contribution}(i) = \frac{\text{factor}(i) \cdot \text{dependent}(i)}{\|\text{factor}\| \cdot \|\text{dependent}\|} \quad \text{for } 1 \leq i \leq n_timepoints$$

Arguments:

- `n_timepoints` integer(int32): Number of timepoints
- `factor` real(real64): Vector of independent variable [n_timepoints]
- `dependent` real(real64): Vector of dependent variable [n_timepoints]
- `mode` integer(int32): Processing mode (1 = Normal, 2 = RAP)
- `contribution` real(real64): Output spike contributions [n_timepoints]
- `ierr` integer(int32): Error code

Listing 61: Spike Contribution Calculation

```
call spike_contribution(factor, dependent, n_timepoints, mode,
    contribution, ierr)
```

(c) **Calculate Contributions (Expert Version)**

Calculates both spike and integrated contributions for a specific factor with pre-allocated work arrays.

Arguments:

- `n_factors` integer(int32): Number of factors in trajectories
- `n_samples` integer(int32): Number of samples in trajectories
- `n_timepoints` integer(int32): Number of timepoints in trajectories
- `trajectories` real(real64): 3D trajectories array [n_factors, n_samples, n_timepoints]
- `i_factor` integer(int32): Factor index to calculate contributions for
- `dependent_idx` integer(int32): Dependent variable index
- `mode` integer(int32): Processing mode (1 = Normal, 2 = RAP)
- `integrated_contribs` real(real64): Output integrated contributions [n_samples]
- `spike_contribs` real(real64): Output spike contributions [n_timepoints, n_samples]
- `temp_factor_vector` real(real64): Work array [n_timepoints]
- `temp_dependent_vector` real(real64): Work array [n_timepoints]
- `ierr` integer(int32): Error code

Listing 62: Expert Contributions Calculation

```
call calc_contributions(trajectories, n_factors, n_samples, &
    n_timepoints, i_factor, dependent_idx, mode, &
    spike_contribs, integrated_contribs, &
    temp_factor_vector, temp_dependent_vector, ierr)
```

(d) **Calculate Contributions (Allocation Version)**

Calculates both spike and integrated contributions with internal memory allocation.

Arguments:

- `n_factors` integer(int32): Number of factors in trajectories
- `n_samples` integer(int32): Number of samples in trajectories
- `n_timepoints` integer(int32): Number of timepoints in trajectories
- `trajectories` real(real64): 3D trajectories array [n_factors, n_samples, n_timepoints]
- `i_factor` integer(int32): Factor index to calculate contributions for
- `dependent_idx` integer(int32): Dependent variable index
- `mode` integer(int32): Processing mode (1 = Normal, 2 = RAP)
- `integrated_contribs` real(real64): Output integrated contributions [n_samples]
- `spike_contribs` real(real64): Output spike contributions [n_timepoints, n_samples]
- `ierr` integer(int32): Error code

Listing 63: Allocation-Based Contributions Calculation

```
call calc_contributions_alloc(trajectories, n_factors, n_samples, &
                             n_timepoints, i_factor, dependent_idx, mode, &
                             spike_contribs, integrated_contribs, ierr)
```

(e) Vector Accessor Routines

Extract specific slices from the 3D trajectories array.

Arguments for Sample Extraction:

- `n_factors` integer(int32): Number of factors
- `n_samples` integer(int32): Number of samples
- `n_timepoints` integer(int32): Number of timepoints
- `trajectories` real(real64): 3D trajectories array
- `i_factor` integer(int32): Factor index
- `i_timepoint` integer(int32): Timepoint index
- `result_vector` real(real64): Output vector [n_samples]
- `ierr` integer(int32): Error code

Listing 64: Extract Samples Vector

```
call get_vec_across_samples(trajectories, n_factors, n_samples, &
                           n_timepoints, i_factor, i_timepoint, &
                           result_vector, ierr)
```

Arguments for Timepoint Extraction:

- `n_factors` integer(int32): Number of factors
- `n_samples` integer(int32): Number of samples
- `n_timepoints` integer(int32): Number of timepoints
- `trajectories` real(real64): 3D trajectories array
- `i_factor` integer(int32): Factor index
- `i_sample` integer(int32): Sample index
- `result_vector` real(real64): Output vector [n_timepoints]
- `ierr` integer(int32): Error code

Listing 65: Extract Timepoints Vector

```
call get_vec_across_timepoints(trajectories, n_factors, n_samples, &
                              n_timepoints, i_factor, i_sample, &
                              result_vector, ierr)
```

(f) Outlier detection

The outlier detection workflow consists of two main pathways: integrated (trajectory-level) analysis and spike (timepoint-level) analysis. You can perform outlier detection by calling the following subroutines in sequence:

i. Calculate Integrated Threshold

The subroutine calculates an empirical threshold for integrated (trajectory-level) contributions using pre-computed permutations. This determines which samples are significant overall across all genes. Call the `calc_integrated_threshold` subroutine with the following arguments:

Input arguments:

- `n_samples` integer(int32): Number of samples
- `contributions` real(real64): 1D array of integrated contributions (size: `n_samples`)
- `percentile_val` real(real64): Percentile value for threshold (0.0–100.0)

- `permutation integer(int32)`: Pre-computed permutation indices for sorted contributions (size: `n_samples`)

Output arguments:

- `threshold real(real64)`: Scalar threshold value
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 66: Integrated Threshold Calculation

```
call calc_integrated_threshold(contributions, n_samples, &
                             percentile_val, threshold, &
                             permutation, ierr)
```

ii. **Detect Integrated Outliers**

The subroutine identifies outlier samples where the integrated contribution exceeds the empirical threshold. Used to find trajectories with significant overall explanatory contributions. Call the `detect_outliers_integrated` subroutine with the following arguments:

Input arguments:

- `n_samples integer(int32)`: Number of samples
- `contributions real(real64)`: Array of integrated contributions (size: `n_samples`)
- `threshold real(real64)`: Scalar threshold value for outlier detection

Output arguments:

- `outlier_mask logical`: 1D logical array indicating outliers (size: `n_samples`)
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 67: Integrated Outlier Detection

```
call detect_outliers_integrated(contributions, n_samples, &
                               threshold, outlier_mask, ierr)
```

iii. **Calculate Spike Thresholds**

The subroutine calculates empirical thresholds for spike contributions using pre-sorted permutations. Computes percentile-based thresholds for each timepoint in spike contribution data. Call the `calc_spike_thresholds` subroutine with the following arguments:

Input arguments:

- `n_timepoints integer(int32)`: Number of timepoints
- `n_samples integer(int32)`: Number of samples
- `spike_contribs real(real64)`: 2D array of spike contributions (size: `n_timepoints` × `n_samples`)
- `percentile_val real(real64)`: Percentile value for threshold (0.0–100.0)
- `permutation integer(int32)`: Pre-computed permutation indices (size: `n_samples` × `n_timepoints`).

Output arguments:

- `thresholds real(real64)`: 1D array of thresholds for each timepoint (size: `n_timepoints`)
- `ierr integer(int32)`: Error code (for error details please check error handling section)

Listing 68: Spike Thresholds Calculation

```
call calc_spike_thresholds(spike_contribs, n_timepoints,
                          n_samples, percentile_val, thresholds, &
                          permutation, ierr)
```

iv. Detect Spike Outliers

The subroutine identifies outliers in spike contribution data by comparing against timepoint-specific thresholds. Performs element-wise comparison across all timepoints and samples. Call the `detect_outliers_spike` subroutine with the following arguments:

Input arguments:

- `n_timepoints` integer(int32): Number of timepoints
- `n_samples` integer(int32): Number of samples
- `spike_contribs` real(real64): 2D array of spike contributions (size: `n_timepoints` × `n_samples`)
- `thresholds` real(real64): Array of thresholds for each timepoint (size: `n_timepoints`)

Output arguments:

- `outlier_mask` logical: 2D logical array indicating outliers (size: `n_timepoints` × `n_samples`)
- `ierr` integer(int32): Error code (for error details please check error handling section)

Listing 69: Spike Outlier Detection

```
call detect_outliers_spike(spike_contribs, n_timepoints, &
                           n_samples, thresholds, outlier_mask, ierr)
```

For convenience, allocation-friendly versions are also available:

- `calc_integrated_threshold_alloc`: Handles internal allocation and sorting
- `calc_spike_thresholds_alloc`: Handles internal allocation and sorting

(g) Process Trajectories with Per-Timepoint Percentiles

Complete pipeline processing with timepoint-specific spike thresholds.

Arguments:

- `n_factors` integer(int32): Number of factors
- `n_samples` integer(int32): Number of samples
- `n_timepoints` integer(int32): Number of timepoints
- `factor_mask_count` integer(int32): Number of true values in factor mask
- `trajectories` real(real64): 3D trajectories array
- `factor_mask` logical: Mask for included factors [`n_factors`]
- `dependent_idx` integer(int32): Dependent variable index
- `mode` integer(int32): Processing mode (1 = Normal, 2 = RAP)
- `percentile` real(real64): Percentile value (0.0–100.0)
- `integrated_contribs` real(real64): Output integrated contributions [`n_samples`, `factor_mask_count`]
- `spike_contribs` real(real64): Output spike contributions [`n_timepoints`, `n_samples`, `factor_mask_count`]
- `thresholds_integrated_contrib` real(real64): Integrated thresholds [`factor_mask_count`]
- `thresholds_spike_contrib` real(real64): Spike thresholds [`n_timepoints`, `factor_mask_count`]
- `outliers_integrated_contrib` logical: Integrated outliers [`n_samples`, `factor_mask_count`]
- `outliers_spike_contrib` logical: Spike outliers [`n_timepoints`, `n_samples`, `factor_mask_count`]
- `ierr` integer(int32): Error code

Listing 70: Complete Pipeline with Per-Timepoint Thresholds

```
call process_trajectories_alloc(trajectories, n_factors, n_samples, &
                               n_timepoints, factor_mask, factor_mask_count, &
                               dependent_idx, mode, percentile, &
```

```

        integrated_contribs, spike_contribs, &
        thresholds_integrated_contrib, &
        outliers_integrated_contrib, &
        thresholds_spike_contrib, &
        outliers_spike_contrib, ierr)

```

(h) Process Trajectories with Global Percentile

Complete pipeline processing with global spike thresholds (flattened approach).

Arguments:

- `n_factors` integer(int32): Number of factors
- `n_samples` integer(int32): Number of samples
- `n_timepoints` integer(int32): Number of timepoints
- `factor_mask_count` integer(int32): Number of true values in factor mask
- `trajectories` real(real64): 3D trajectories array
- `factor_mask` logical: Mask for included factors [`n_factors`]
- `dependent_idx` integer(int32): Dependent variable index
- `mode` integer(int32): Processing mode (1 = Normal, 2 = RAP)
- `percentile` real(real64): Percentile value (0.0–100.0)
- `integrated_contribs` real(real64): Output integrated contributions [`n_samples`, `factor_mask_count`]
- `spike_contribs` real(real64): Output spike contributions [`n_timepoints`, `n_samples`, `factor_mask_count`]
- `thresholds_integrated_contrib` real(real64): Integrated thresholds [`factor_mask_count`]
- `thresholds_spike_contrib` real(real64): Spike thresholds [`factor_mask_count`] (1D)
- `outliers_integrated_contrib` logical: Integrated outliers [`n_samples`, `factor_mask_count`]
- `outliers_spike_contrib` logical: Spike outliers [`n_timepoints`, `n_samples`, `factor_mask_count`]
- `ierr` integer(int32): Error code

Listing 71: Complete Pipeline with Global Threshold

```

call process_trajectories_flat_alloc(trajectories, n_factors, &
                                     n_samples, n_timepoints, factor_mask, &
                                     factor_mask_count, dependent_idx, mode, &
                                     percentile, integrated_contribs, &
                                     spike_contribs, thresholds_integrated_contrib, &
                                     outliers_integrated_contrib, &
                                     thresholds_spike_contrib, &
                                     outliers_spike_contrib, ierr)

```

Processing Modes:

- `MODE_NORMAL` = 1: Returns cosine similarity for normal vectors
- `MODE_RAP` = 2: Returns arccos of cosine similarity for RAP projected vectors

Key Features:

- **Pure Subroutines:** Most routines are pure with no side effects
- **Robust Error Handling:** Comprehensive error checking and validation
- **Memory Efficient:** Expert versions allow pre-allocated work arrays
- **Mathematically Sound:** Uses cosine similarity with proper normalization
- **Flexible Processing:** Supports both per-timepoint and global thresholding

3.1.8 Fortran Utils

1. Binary Search Trees for 1D Range Queries

Binary Search Trees provide efficient indexing and range query capabilities for one-dimensional data using flat array implementations for optimal memory performance.

Fortran BST Implementation

(a) Build BST Index

Constructs a Binary Search Tree index by sorting the input values and storing permutation indices. Uses an efficient stack-based sorting algorithm.

Input:

- `values` `real(real64)`: 1D array of values to index
- `num_values` `integer(int32)`: Number of elements

Output:

- `sorted_indices` `integer(int32)`: Permutation indices after sorting (1-based)
- `left_stack`, `right_stack` `integer(int32)`: Sorting workspace
- `ierr` `integer(int32)`: Error code

```
call build_bst_index(values, num_values, sorted_indices, left_stack,  
                    right_stack, ierr)
```

(b) Get Sorted Value

Retrieves the value at a specific position in the sorted order defined by the BST index.

Input:

- `values` `real(real64)`: Original values array
- `sorted_indices` `integer(int32)`: BST index array (1-based)
- `position` `integer(int32)`: Position in sorted order (1-based)

Output:

- `sorted_value` `real(real64)`: Value at specified position
- `ierr` `integer(int32)`: Error code

```
sorted_value = get_sorted_value(values, sorted_indices, position, ierr  
                                )
```

(c) BST Range Query

Efficiently finds all values within a specified range using the precomputed BST index.

Input:

- `values` `real(real64)`: Original values array
- `sorted_indices` `integer(int32)`: BST index array (1-based)
- `num_values` `integer(int32)`: Number of elements
- `lower_bound`, `upper_bound` `real(real64)`: Range boundaries

Output:

- `output_indices` `integer(int32)`: Indices of matching values (1-based)
- `num_matches` `integer(int32)`: Number of matches found
- `ierr` `integer(int32)`: Error code

```
call bst_range_query(values, sorted_indices, num_values, &  
                    lower_bound, upper_bound, output_indices, &  
                    num_matches, ierr)
```

2. k-d Trees for Multi-Dimensional Indexing

k-d trees provide efficient space partitioning for multi-dimensional data with support for both Euclidean and spherical geometries.

Fortran k-d Tree Implementation

(a) Build k-d Tree Index

Constructs a k-d tree index using stack-based non-recursive partitioning along specified dimensions.

Input:

- `points` `real(real64)`: Data points matrix (dimensions $\times$ points)
- `num_dimensions` `integer(int32)`: Number of dimensions
- `num_points` `integer(int32)`: Number of points
- `dimension_order` `integer(int32)`: Dimension ordering for splitting (1-based)

Output:

- `kd_indices` `integer(int32)`: Output index array (1-based)
- `workspace` `integer(int32)`: Workspace array
- `value_buffer` `real(real64)`: Value buffer for sorting
- `permutation` `integer(int32)`: Permutation array
- `left_stack`, `right_stack` `integer(int32)`: Sorting stacks
- `recursion_stack` `integer(int32)`: Recursion stack
- `ierr` `integer(int32)`: Error code

```
call build_kd_index(points, num_dimensions, num_points, kd_indices, &
                   dimension_order, workspace, value_buffer, &
                   permutation, left_stack, right_stack, &
                   recursion_stack, ierr)
```

(b) Build Spherical k-d Tree

Specialized k-d tree construction optimized for unit vectors on hyperspheres.

Input:

- `vectors` `real(real64)`: Unit vectors matrix (dimensions $\times$ vectors)
- `num_dimensions` `integer(int32)`: Number of dimensions
- `num_vectors` `integer(int32)`: Number of vectors
- `dimension_order` `integer(int32)`: Dimension ordering for splitting (1-based)

Output:

- `sphere_indices` `integer(int32)`: Output index array (1-based)
- Various workspace arrays (same as `build_kd_index`)
- `ierr` `integer(int32)`: Error code

```
call build_spherical_kd(vectors, num_dimensions, num_vectors, &
                       sphere_indices, dimension_order, workspace, &
                       value_buffer, permutation, left_stack, &
                       right_stack, recursion_stack, ierr)
```

(c) Get Point from k-d Index

Retrieves point coordinates from specified position in k-d tree ordering.

Input:

- `points` `real(real64)`: Original points matrix
- `kd_indices` `integer(int32)`: k-d tree index array (1-based)
- `position` `integer(int32)`: Position in index (1-based)

Output:

- `point_values` `real(real64)`: Retrieved point values
- `ierr` `integer(int32)`: Error code

```
call get_kd_point(points, kd_indices, position, point_values, ierr)
```


3.2 Python Pipeline

The Python workflow is exposed through the `tensoromics_functions.py` module. The functions provide a high-level interface that mirrors the Fortran pipeline.

3.2.1 Data Management

1. Data Reading and Processing

• Read Gene IDs from TSV File

The function reads gene identifiers from a tab-separated value (TSV) file, extracting specific gene IDs from a designated column while skipping header rows. This function is typically used as the first step in processing gene expression data to obtain the complete list of gene identifiers. Call the `read_gene_ids_from_tsv_file` function with the following arguments:

- `filename`: Path to the TSV file containing gene IDs (string or list with single file)
- `n_genes`: Total number of genes in the file
- `gene_ids_len`: Maximum character length for each gene ID
- `n_header_rows`: Number of header rows to skip at the beginning of the file
- `gene_col`: Column index (1-based) where gene IDs are located

Returns:

- `gene_ids`: Numpy array of shape `(n_genes)` containing gene ID strings

Listing 72: Python Read Gene IDs from TSV File

```
gene_ids = read_gene_ids_from_tsv_file(filename, n_genes, gene_ids_len,
                                       n_header_rows, gene_col)
```

• Read Expression Vectors

The function reads gene expression data from multiple tabular files (CSV/TSV format), extracting expression values for specified genes across multiple samples. This function consolidates expression data from multiple files into a single 2D array. Call the `read_expression_vectors` function with the following arguments:

- `file_list`: List of file paths containing expression data
- `gene_ids`: Array of gene IDs to extract expression values for
- `n_samples`: Number of samples
- `n_header_rows`: Number of header rows to skip in each file
- `gene_col`: Column index (1-based) containing gene identifiers
- `value_cols`: List of column indices (1-based) containing expression values
- **(Optional)** `delimiter`: Delimiter used in the files (default: tab)

Returns:

- `expression_vectors`: 2D numpy array of shape `(n_samples, n_genes)` containing expression values as 64-bit floats

Listing 73: Python Read Expression Vectors

```
expression_vectors = read_expression_vectors(file_list, gene_ids,
                                             n_samples, n_header_rows, gene_col, value_cols, delimiter='\t')
```

• Read OrthoFinder File

The function processes an OrthoFinder output file to map genes to their respective orthologous families. This function establishes the gene-family relationships essential for downstream orthology-based analyses. Call the `read_orthofinder_file` function with the following arguments:

- **filename**: Path to the OrthoFinder TSV file (string or list with single file)
- **gene_ids**: Complete list of gene identifiers to map
- **family_ids_len**: Maximum character length for family IDs
- **n_families**: Total number of gene families expected

Return dictionary:

- **family_ids**: Numpy array of shape (**n_families**,) containing family ID strings
- **gene_to_fam**: Numpy array of shape (**n_genes**,) containing family indices for each gene (0 = unassigned)

Listing 74: Python Read OrthoFinder File

```
family_result = read_orthofinder_file(filename, gene_ids,
                                     family_ids_len, n_families)
```

Complete Data Loading Workflow Example:

Listing 75: Python Complete Data Loading Workflow

```
# Step 1: Load gene identifiers
gene_ids = read_gene_ids_from_tsv_file(
    filename="data/expression_data.tsv",
    n_genes=50000,
    gene_ids_len=32,
    n_header_rows=1,
    gene_col=1
)

# Step 2: Load expression data from multiple samples
samples = ["sample1.tsv", "sample2.tsv", "sample3.tsv", "sample4.tsv"]
expression_vectors = read_expression_vectors(
    file_list=samples,
    gene_ids=gene_ids,
    n_samples=12, # assuming 3 samples per file
    n_header_rows=1,
    gene_col=1,
    value_cols=[2, 3]
)

# Step 3: Load orthology information
orthology_data = read_orthofinder_file(
    filename="data/Orthogroups.tsv",
    gene_ids=gene_ids,
    family_ids_len=32,
    n_families=15000
)
```

2. Data Filtering and Validation

• Filter Unassigned Genes

The function filters out genes that are not assigned to any family (where `gene_to_fam = 0`). This is essential for cleaning data before orthology-based analyses, ensuring only genes with valid family assignments are processed. Call the `filter_unassigned_genes` function with the following arguments:

- **gene_to_fam**: Family assignment array where 0 indicates unassigned genes

Return dictionary:

- **mask**: Boolean array indicating which genes to keep (1 = assigned to family, 0 = unassigned)
- **n_genes_kept**: Number of genes kept after filtering

Listing 76: Python Filter Unassigned Genes

```
filtered_genes = filter_unassigned_genes(gene_to_fam)
```

- **Validate Gene to Family Mapping**

The function validates the gene-to-family mapping array, ensuring all family indices are within valid range and properly assigned. This prevents downstream errors from invalid family indices. Call the `validate_gene_to_family_mapping` function with the following arguments:

- **gene_to_fam**: Array mapping each gene to a family index (0 = unassigned)
- **n_families**: Total number of families

Listing 77: Python Validate Gene to Family Mapping

```
validate_gene_to_family_mapping(gene_to_fam, n_families)
```

- **Validate Expression Data**

The function validates expression data, checking for NaN/Inf values and optionally verifying non-negative values. This ensures data quality before computational analyses. Call the `validate_expression_data` function with the following arguments:

- **expression_vectors**: 2D array of expression data (n_samples x n_genes)
- **(Optional) check_non_negative**: Whether to check for non-negative values (default: True)

Listing 78: Python Validate Expression Data

```
validate_expression_data(expression_vectors, check_non_negative=True)
```

- **Validate Family Centroids**

The function validates family centroids data, checking for NaN/Inf values and proper data structure. Centroids represent the central tendency of each gene family’s expression profile. Call the `validate_family_centroids` function with the following arguments:

- **family_centroids**: 2D array of family centroids (samples x families)

Listing 79: Python Validate Family Centroids

```
validate_family_centroids(family_centroids)
```

- **Validate Shift Vectors**

The function validates shift vectors, ensuring data types are correct and structure matches expected patterns. Shift vectors represent the deviation of each gene from its family centroid. Call the `validate_shift_vectors` function with the following arguments:

- **shift_vectors**: 2D array of shift vectors (samples × genes)
- **expression_vectors**: 2D array of expression data (samples × genes)
- **family_centroids**: 2D array of family centroids (samples × families)
- **gene_to_fam**: Array mapping each gene to a family index
- **n_genes**: Number of genes
- **n_samples**: Number of samples
- **n_families**: Number of families

Listing 80: Python Validate Shift Vectors

```
validate_shift_vectors(shift_vectors, expression_vectors,
                      family_centroids, gene_to_fam, n_genes, n_samples, n_families)
```

- **Validate String Array Uniqueness**

The function validates that all strings in an array are unique, using a hashset internally. Note:

This may increase memory usage temporarily for large datasets. Call the `validate_string_array_uniqueness` function with the following arguments:

- `strings`: Array of strings to validate (typically gene IDs or family IDs)

Listing 81: Python Validate String Array Uniqueness

```
validate_string_array_uniqueness(strings)
```

- **Validate Data Structure**

The function performs comprehensive validation of the overall data structure consistency, confirming sizes and dependencies between different data components as far as possible. Call the `validate_data_structure` function with the following arguments:

- `n_genes`: Number of genes
- `n_families`: Number of families
- `n_samples`: Number of samples
- `gene_ids`: Array of gene IDs
- `gene_family_ids`: Array of family IDs
- `gene_to_fam`: Array mapping each gene to a family index
- `expression_vectors`: 2D array of expression data (samples $\times$ genes)
- `family_centroids`: 2D array of family centroids (samples $\times$ families)
- `shift_vectors`: 2D array of shift vectors (genes $\times$ samples)

Listing 82: Python Validate Data Structure

```
validate_data_structure(n_genes, n_families, n_samples, gene_ids,
                      gene_family_ids, gene_to_fam, expression_vectors, family_centroids,
                      shift_vectors)
```

- **Validate All Data**

The function performs comprehensive validation of all data components in one call. This function executes all individual validations sequentially for complete data quality assurance. Call the `validate_all_data` function with the following arguments:

- `n_genes`: Number of genes
- `n_families`: Number of families
- `n_samples`: Number of samples
- `gene_ids`: Array of gene IDs
- `gene_family_ids`: Array of family IDs
- `gene_to_fam`: Array mapping each gene to a family index
- `expression_vectors`: 2D array of expression data (samples $\times$ genes)
- `family_centroids`: 2D array of family centroids (samples $\times$ families)
- `shift_vectors`: 2D array of shift vectors (samples $\times$ genes)

Listing 83: Python Validate All Data

```
validate_all_data(n_genes, n_families, n_samples, gene_ids,
                 gene_family_ids, gene_to_fam, expression_vectors, family_centroids,
                 shift_vectors)
```

Data Filtering and Validation Workflow Example:

Listing 84: Python Data Filtering and Validation Workflow

```
# After loading data, filter unassigned genes
filter_result = filter_unassigned_genes(gene_to_fam)
mask = np.array(filter_result['mask'], dtype=bool)
n_genes_kept = filter_result['n_genes_kept']

# Apply filter to all gene-based arrays
filtered_gene_ids = gene_ids[mask]
filtered_expression = expression_vectors[:, mask]
filtered_gene_to_fam = gene_to_fam[mask]

# Run comprehensive validations
validate_string_array_uniqueness(filtered_gene_ids)
validate_string_array_uniqueness(family_ids)
validate_expression_data(filtered_expression, check_non_negative=True)
validate_gene_to_family_mapping(filtered_gene_to_fam, n_families)

# For complete analysis with centroids and shift vectors
validate_data_structure(
    n_genes_kept, n_families, n_samples,
    filtered_gene_ids, family_ids, filtered_gene_to_fam,
    filtered_expression, family_centroids, shift_vectors
)

# Or use the comprehensive validation
validate_all_data(
    n_genes_kept, n_families, n_samples,
    filtered_gene_ids, family_ids, filtered_gene_to_fam,
    filtered_expression, family_centroids, shift_vectors
)
```

3.2.2 Data Persistence

1. Python Interface for Array Serialization and Deserialization

The Python interface provides bindings to the Fortran serialization and deserialization routines, enabling efficient handling of multi-dimensional arrays. This interface supports NumPy arrays with native data type compatibility and Fortran-style column-major memory layout.

(a) General Architecture

The Python module `tensoromics_functions` provides ten main functions for array handling:

- `tox_serialize_int_nd(array, filename)` - Serialize N-dimensional integer arrays.
- `tox_deserialize_int_nd(filename)` - Deserialize N-dimensional integer arrays.
- `tox_serialize_real_nd(array, filename)` - Serialize N-dimensional real arrays.
- `tox_deserialize_real_nd(filename)` - Deserialize N-dimensional real arrays.
- `tox_serialize_char_nd(array, filename)` - Serialize N-dimensional character arrays.
- `tox_deserialize_char_nd(filename)` - Deserialize N-dimensional character arrays.
- `tox_serialize_logical_nd(array, filename)` - Serialize N-dimensional logical arrays.
- `tox_deserialize_logical_nd(filename)` - Deserialize N-dimensional logical arrays.
- `tox_serialize_complex_nd(array, filename)` - Serialize N-dimensional complex arrays.
- `tox_deserialize_complex_nd(filename)` - Deserialize N-dimensional complex arrays.

- `get_array_metadata(filename)` - Retrieve array dimensions and metadata
- (b) **Serialization Functions:**
- **array** (input): NumPy array to be serialized
Type: `numpy.ndarray` with dtype:
 - `np.int32` for integer
 - `np.float64` for float values
 - `U5` for characters
 - `boolean` for logical
 - `np.complex128` for complex numbers**Layout:** Column-major (Fortran) order (`order='F'`)
Dimensions: Since the array is being passed flat, theoretically all dimensions are supported, but practical use cases are typically 1D to 5D
 - **filename** (input): Name of the binary file to create
Type: `str`

Listing 85: Example usage of serialization function

```
tox_serialize_int_nd(int_array, filename)
tox_serialize_real_nd(real_array, filename)
```

- (c) **Deserialization Functions:**
- **filename** (input): Name of the binary file to read
Type: `str`
 - **Return Value:** Deserialized NumPy array
Type: `numpy.ndarray` with original dtype and dimensions
Layout: Column-major (Fortran) order

Listing 86: Example usage of deserialization function

```
restored_int = tox_deserialize_int_nd(filename)
restored_real = tox_deserialize_real_nd(filename)
```

- (d) **Implementation Details**
- The Python interface uses Fortran-style column-major array ordering for optimal compatibility
 - All arrays are converted to the appropriate Fortran data types before serialization
 - The interface automatically handles metadata including dimensions and data types
 - Error handling is implemented through Python exceptions
- (e) **Usage Example**

Listing 87: Full example of serialization and deserialization

```
# Real array example
real_array = np.array([1.5, 2.3, 3.2, 4.0, 5.0], dtype=np.float64,
                      order='F')
tox_serialize_real_nd(real_array, "test_real_1d.bin")
restored_real = tox_deserialize_real_nd("test_real_1d.bin")

# Character array example with metadata
char_array = np.asfortranarray(["hello", "world"], dtype='U5')
tox_serialize_char_nd(char_array, "test_char_1d.bin")

restored_char = tox_deserialize_char_nd("test_char_1d.bin")
```

```
# Multi-dimensional examples
int_2d = np.array([[1, 2, 3], [4, 5, 6]], dtype=np.int32, order='F')
tox_serialize_int_nd(int_2d, "test_int_2d.bin")

restored_2d = tox_deserialize_int_nd("test_int_2d.bin")
```

(f) **Compatibility Notes**

- Requires NumPy installation
- Arrays must use Fortran-style column-major ordering
- Character arrays must use fixed-length string data types
- The interface supports N dimensions, but practical use cases are typically 1D to 5D
- Data types are restricted to int32, float64, and U5 for characters

2. Standard Archive Operations

• Create Zip Archive (Low-Level)

The low-level function creates a zip archive from keys and filenames, directly calling the Fortran backend. Use this function for custom data. Note that the identifiers at position i belongs to position i of the filenames. All files named in the filenames have to be saved to the disk before the function is called. For more details on how to save arrays, see section array serialization and deserialization. Call the `create_zip_archive` function with the following arguments:

- `zip_filename`: Name of the zip file to create
- `keys`: List of keys for the manifest (identifiers for each file)
- `filenames`: List of filenames to include in archive

Listing 88: Python Create Zip Archive (Low-Level)

```
create_zip_archive(zip_filename, keys, filenames)
```

• Extract Zip Archive

The function extracts a zip archive created by `create_zip_archive` and returns a dictionary mapping data keys to extracted filenames. All extracted files will be written to the disk. For more details on how to continue working with the files, see section array serialization and deserialization. Call the `extract_zip_archive` function with the following arguments:

- `zip_filename`: Path to the zip file to extract

Returns:

- Dictionary mapping data keys to extracted temporary filenames

Listing 89: Python Extract Zip Archive

```
extract_zip_archive(zip_filename)
```

• Save TOX Data (High-Level)

The high-level function saves TOX data to a zip archive, handling serialization, and archive creation for standard-conform tox gene data. This function automatically manages temporary files and cleanup. Call the `save_tox_data` function with the following arguments:

- `zip_filename`: Name of the zip file to create
- **(Optional)** `gene_ids`: Array of gene IDs
- **(Optional)** `expression_vectors`: 2D array of expression data
- **(Optional)** `gene_to_fam`: Array mapping genes to families
- **(Optional)** `family_ids`: Array of family IDs

- **(Optional)** `family_centroids`: 2D array of family centroids
- **(Optional)** `shift_vectors`: 2D array of shift vectors
- Corresponding `*_name` parameters for temporary filenames (e.g., `gene_ids_name`, `expression_vectors_name` etc.)

Listing 90: Python Save TOX Data

```
save_tox_data(zip_filename, gene_ids=None, expression_vectors=None,
              gene_to_fam=None, family_ids=None, family_centroids=None,
              shift_vectors=None, gene_ids_name=None, expression_vectors_name=
              None, gene_to_fam_name=None, family_ids_name=None,
              family_centroids_name=None, shift_vectors_name=None)
```

• Read TOX Data

The function reads data from a zip archive created by `save_tox_data`, extracting and deserializing the requested data components for standard-conform tox gene data. This function automatically handles file extraction, cleanup and array allocation. Call the `read_tox_data` function with the following arguments:

- `zip_filename`: Name of the zip file to read
- **(Optional)** `load_gene_ids`: Whether to load gene IDs (default: False)
- **(Optional)** `load_expression_vectors`: Whether to load expression vectors (default: False)
- **(Optional)** `load_gene_to_fam`: Whether to load gene-to-family mapping (default: False)
- **(Optional)** `load_family_ids`: Whether to load family IDs (default: False)
- **(Optional)** `load_family_centroids`: Whether to load family centroids (default: False)
- **(Optional)** `load_shift_vectors`: Whether to load shift vectors (default: False)

Return dictionary:

- `gene_ids`: Loaded gene IDs array (if requested)
- `expression_vectors`: Loaded expression vectors array (if requested)
- `gene_to_fam`: Loaded gene-to-family mapping array (if requested)
- `family_ids`: Loaded family IDs array (if requested)
- `family_centroids`: Loaded family centroids array (if requested)
- `shift_vectors`: Loaded shift vectors array (if requested)

Listing 91: Python Read TOX Data

```
read_tox_data(zip_filename, load_gene_ids=False,
              load_expression_vectors=False, load_gene_to_fam=False,
              load_family_ids=False, load_family_centroids=False,
              load_shift_vectors=False)
```

Data Archiving Workflow Examples:

Listing 92: Python Standard TOX Data Archiving

```
# Save complete TOX dataset
save_tox_data(
    zip_filename="complete_tox_dataset.zip",
    gene_ids=filtered_gene_ids,
    gene_ids_name="gene_ids.bin",
    expression_vectors=filtered_expression,
    expression_vectors_name="expression.bin",
    gene_to_fam=filtered_gene_to_fam,
    gene_to_fam_name="gene_to_fam.bin",
```



```

        family_ids=family_ids,
        family_ids_name="family_ids.bin",
        family_centroids=family_centroids,
        family_centroids_name="centroids.bin",
        shift_vectors=shift_vectors,
        shift_vectors_name="shift_vectors.bin"
    )

    # Save partial dataset
    save_tox_data(
        zip_filename="basic_data.zip",
        gene_ids=filtered_gene_ids,
        gene_ids_name="gene_ids.bin",
        expression_vectors=filtered_expression,
        expression_vectors_name="expression.bin"
    )

    # Read data back with selective loading
    loaded_data = read_tox_data(
        zip_filename="complete_tox_dataset.zip",
        load_gene_ids=True,
        load_expression_vectors=True,
        load_gene_to_fam=True
    )

```

Listing 93: Python Custom Archive Creation and Extraction

```

# Create custom archive with non-standard data
create_zip_archive(
    zip_filename="custom_data.zip",
    keys=["3d_tensor", "results"],
    filenames=["tensor_data.bin", "analysis_results.bin"]
)

# Extract and access custom archive
file_mapping = extract_zip_archive("custom_data.zip")
print("Files in archive:", file_mapping)
# Output: {'3d_tensor': 'tensor_data.bin', 'results': '
    analysis_results.bin'}

```

Listing 94: Python Mixed Standard and Custom Data Archiving

```

# Serialize custom arrays
tox_serialize_real_nd(custom_array, "custom_data.bin")
tox_serialize_int_nd(another_array, "another_data.bin")

# Create mixed archive with both standard and custom data
create_zip_archive(
    zip_filename="mixed_archive.zip",
    keys=["standard_expression", "custom_feature", "analysis_results"
    ],
    filenames=["expression.bin", "custom_data.bin", "another_data.bin"
    ]
)

```

```
# Later, extract and process
mapping = extract_zip_archive("mixed_archive.zip")
expression_data = tox_deserialize_real_nd(mapping["standard_expression"])
custom_features = tox_deserialize_real_nd(mapping["custom_feature"])
results = tox_deserialize_int_nd(mapping["analysis_results"])
```

3.2.3 Data Processing

1. Normalization of Expression Values

The function runs a complete normalization pipeline on gene expression data. It first applies standard deviation normalization, then quantile normalization, followed by replicate averaging, and finally a log2 transformation. The result is returned in an array, which contains the fully normalized expression values. Call the `tox_normalization_pipeline` function with the following arguments:

- **input_matrix**: Numeric expression vector matrix with size `genes × tissues`
- **group_starts**: Integer array, start column index for each replicate group (1-based)
- **group_counts**: Integer array, number of columns per replicate group

Return:

- Numpy array with $\log_2(x+1)$ normalized expression with size `(genes × groups)`

Listing 95: Python Normalization Pipeline

```
tox_normalization_pipeline(input_matrix, group_starts, group_counts)
```

Alternatively you can calculate each step of the normalization pipeline manually by calling the following functions in sequence:

(a) Normalize by Standard Deviation

The function normalizes each gene's expression vector using `sqrt(mean(x2))` across tissues (not classical standard deviation). Call the `tox_normalize_by_std_dev` function with the following arguments:

- **input_matrix**: A numeric matrix with genes as rows and tissues as columns

Returns:

- Numpy array: Contains a normalized matrix with same dimensions as input

Listing 96: Python Standard Deviation Normalization

```
tox_normalize_by_std_dev(input_matrix)
```

(b) Quantile Normalization

The function performs quantile normalization of a gene expression matrix (F42-compliant) and computes average expression per rank across tissues. Call the `tox_quantile_normalization` function with the following arguments:

- **input_matrix**: A numeric matrix with genes as rows and tissues as columns

Returns:

- Numpy array: Quantile-normalized matrix with same dimensions as input

Listing 97: Python Quantile Normalization

```
tox_quantile_normalization(input_matrix)
```

(c) **Calculate Tissue Averages**

The function calculates tissue averages by averaging replicates within each group. For each group of tissue replicates, it computes the average expression per gene. The input matrix is column-major, flattened as a 1D array. Call the `tox_calculate_tissue_averages` function with the following arguments:

- `input_matrix`: A numeric matrix with genes as rows and tissue replicates as columns
- `group_starts`: Array of starting column indices for each group (1-based for Fortran)
- `group_counts`: Array of counts for each group

Returns:

- Numpy array: Matrix with genes as rows and averaged tissues as columns

Listing 98: Python Tissue Averages Calculation

```
tox_calculate_tissue_averages(input_matrix, group_starts, group_counts
)
```

(d) **Log2 Transformation**

The function applies $\log_2(x + 1)$ transformation to each element of the input matrix. This is computed via $\log(x + 1) / \log(2)$ for compatibility with WebAssembly (WASM). Call the `tox_log2_transformation` function with the following arguments:

- `input_matrix`: A numeric matrix with genes as rows and tissues as columns

Returns:

- Numpy array: Log2-transformed matrix with same dimensions as input

Listing 99: Python Log2 Transformation

```
tox_log2_transformation(input_matrix)
```

2. Calculate Fold Change

If fold change calculation is needed, this function can be used after normalization. The function calculates $\log_2$ fold changes between condition and control columns. For each control-condition pair, this function computes the $\log_2$ fold change by subtracting the expression value in the control column from the corresponding value in the condition column, for all genes. Call the `tox_calculate_fold_changes` function with the following arguments:

- `input_matrix`: A numeric matrix with genes as rows and tissues/conditions as columns
- `control_cols`: Array of control column indices (1-based for Fortran)
- `condition_cols`: Array of condition column indices (1-based for Fortran)

Returns:

- Numpy array: Matrix with genes as rows and fold change values as columns

Listing 100: Python Fold Change Calculation

```
tox_calculate_fold_changes(input_matrix, control_cols, condition_cols)
```

3.2.4 Family-Based Analysis

1. Calculation of family Centroids

The function computes the centroid (mean expression vector) for each gene family in a gene expression dataset. It iterates over all families, selects the relevant genes for each family and calculates the mean vector for each family. There are two possible modes to use. To use all genes for calculation, select `all` mode. If you want to specify relevant genes for calculation,

select `orthologs` mode. It is important that you provide a logical `ortholog_set` array when using `orthologs` mode. Call the `tox_group_centroid` function with the following arguments:

- `expression_vectors`: Expression vector matrix with size `n_axes × n_genes`
- `gene_to_family`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)
- `n_families`: Total number of gene families
- `mode`: Calculation mode ("all" or "orthologs")
- (Optional) `ortholog_set`: Logical array indicating if a gene is part of the calculation subset (only used in `orthologs` mode)

Return:

- `centroids_out`: Matrix of calculated family centroid vectors with size `n_axes × n_families` (read-only)

Listing 101: Python Calculation of Family Centroids

```
tox_group_centroid(expression_vectors, gene_to_family, n_families, mode,
                   ortholog_set)
```

(a) **Calculate Mean Vector**

Alternatively you can calculate a single mean vector for a given set of vectors by calling the `mean_vector` function directly with the following arguments:

- `expression_vectors` `real(real64)`: Expression vector matrix with size `n_axes × n_genes`
- `gene_indices` `integer(int32)`: An array containing the column indices of the selected genes in `expression_vectors`

Return:

- `centroid`: Calculated mean vector (centroid) with length of `n_axes` (read-only)

Listing 102: Python Mean Vector Calculation

```
tox_mean_vector(expression_vectors, gene_indices)
```

2. Calculate Distance to Centroid

For each gene in the gene expression matrix, the function computes the Euclidean distance to its corresponding family centroid and returns a read-only distances array containing the distance value for each expression vector to its family centroid. Call the `tox_distance_to_centroid` function with the following arguments:

- `genes`: Gene expression matrix array with size of `(n_genes × d)`
- `centroids`: Family centroid matrix array with size of `(n_families × d)`
- `gene_to_fam`: Gene-to-family mapping array with the length of `n_genes` (0 = no family, >0 = family index)
- `d`: Gene expression matrix dimension as number

Listing 103: Python Distance to Centroid Calculation

```
tox_distance_to_centroid(genes, centroids, gene_to_fam, d)
```

(a) **Calculate Euclidean Distance**

Alternatively you can calculate a single Euclidean distance between two vectors manually by calling `tox_euclidean_distance` function with the following arguments:

- `vec1`: First vector as array
- `vec2`: Second vector as array

Both vectors must have same length!

Listing 104: Python Euclidean distance calculation

```
tox_euclidean_distance(vec1, vec2)
```

3. Detect Gene Outliers

The function first computes family scaling, then calculates relative distance indices (RDI) and identifies outliers based on the specified percentile threshold. It returns a read-only dictionary containing the outlier boolean values and intermediate results. Call the `tox_detect_outliers` function with the following arguments:

- **distances:** Array of distances for each gene to its centroid
- **gene_to_fam:** Gene-to-family mapping array with the same length as the distances array (0 = no family, >0 = family index)
- **(Optional) percentile:** Percentile threshold value for outlier detection (default is 95)

Return dictionary:

- **outliers:** Boolean array indicating whether each gene is an outlier
- **loess_x:** Array of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y:** Array of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **loess_n:** Array with numbers of genes in each family used to compute the corresponding median and standard deviation

Listing 105: Python Gene Outlier Calculation

```
tox_detect_outliers(distances, gene_to_fam, percentile=95.0)
```

Alternatively you can calculate each step of `tox_detect_outliers` manually with the following functions:

(a) Compute Family Scaling

The function calculates a scaling factor for each gene family by computing the median and standard deviation of gene-to-centroid distances within each family, then applies LOESS smoothing to these values. Call the `tox_compute_family_scaling` function with the following arguments:

- **distances:** Array of distances for each gene to its centroid
- **gene_to_fam:** Gene-to-family mapping array with the same length as the distances array (0 = no family, >0 = family index)

Return dictionary:

- **dscale:** Array of scaling factors per family
- **loess_x:** Array of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y:** Array of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **indices_used:** Array of indices of families that have been included in the calculation

Listing 106: Python Compute Family Scaling

```
tox_compute_family_scaling(distances, gene_to_fam)
```

(b) Compute Family Scaling (Expert Version)

This function is the expert version of `tox_detect_outliers` and requires pre-allocated work arrays for maximum performance and control. It can be used when you need fine-grained control over memory allocation or you are calling this function many times in a tight loop. Call the `tox_compute_family_scaling_expert` function with the following arguments:

- **distances**: Array of distances for each gene to its centroid
- **gene_to_fam**: Gene-to-family mapping array with the same length as the distances array (0 = no family, >0 = family index)
- **perm_tmp**: Permutation array for sorting with the same length as the distances array
- **stack_left_tmp**: Stack array for sorting with the same length as the distances array
- **stack_right_tmp**: Stack array for sorting with the same length as the distances array
- **family_distances**: Work array for sorting with the same length as the distances array

Return dictionary:

- **dscale**: Array of scaling factors per family
- **loess_x**: Array of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y**: Array of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **indices_used**: Array of indices of families that have been included in the calculation
- **perm_tmp**: Permutation array for sorting
- **stack_left_tmp**: Stack array for sorting
- **stack_right_tmp**: Stack array for sorting
- **family_distances**: Work array for sorting

Listing 107: Python Compute Family Scaling (Expert Version)

```
tox_compute_family_scaling_expert(distances, gene_to_fam, perm_tmp,
    stack_left_tmp, stack_right_tmp, family_distances)
```

(c) **Compute Relative Distance Index (RDI)**

The function calculates the Relative Distance Index (RDI) for each gene by dividing its Euclidean distance to the family centroid by the scaling factor of its family. The function returns a numpy array with RDI values for each gene. Call the `tox_compute_rdi` function with the following arguments:

- **distances**: Array of distances for each gene to its centroid
- **gene_to_fam**: Gene-to-family mapping array with the same length as the distances array (0 = no family, >0 = family index)
- **dscale**: Array of scaling factors per family

Listing 108: Python Compute Relative Distance Index (RDI)

```
tox_compute_rdi(distances, gene_to_fam, dscale)
```

(d) **Identify Outliers**

The function identifies outliers based on the top percentile of RDI values. Call the `tox_identify_outliers` function with the following arguments:

- **rdi**: Array of RDI values for each gene
- **(Optional) threshold**: Fixed RDI threshold number (if None, percentile value is used)
- **(Optional) percentile**: Percentile threshold number for outlier detection (default is 95)

Return dictionary:

- **outliers**: Boolean array indicating whether each gene is an outlier
- **threshold**: Actual threshold value used for outlier detection

Listing 109: Python Identify Outliers

```
tox_identify_outliers(rdi, threshold=None, percentile=95.0)
```

4. Compute Normalized Tissue Versatility:

The function calculates a normalized tissue versatility score for selected gene expression vectors. It is based on the angle between each gene expression vector and the space diagonal. Versatility is normalized to [0, 1], where 0 means uniform expression 1 means expression in only one axis. Call the `tox_compute_tissue_versatility` function with the following parameters:

- **expression_vectors**: Expression vector matrix array with size `n_axes × n_vectors`
- **vector_selection**: Logical array with length of `n_vectors` to specify which vectors should be included in the calculation
- **axes_selection**: Logical array with length of `n_axes` to specify which axes should be included in the calculation

Return dictionary:

- **tissue_versatilities**: Array containing the calculated tissue versatilities for selected vectors
- **tissue_angles_deg**: Array containing the calculated angles in degrees for selected vectors
- **n_selected_vectors**: Number of vectors processed
- **n_selected_axes**: Number of axes used in calculation

Listing 110: Python Tissue Versatility Calculation

```
tox_calculate_tissue_versatility(expression_vectors, vector_selection, axis_selection)
```

3.2.5 Relative Axis Plane Projection

Placeholder

3.2.6 Contribution Analysis

1. Fortran Core Contribution Analysis Functions

2. Trajectory Contribution Calculation

Computes integrated trajectory-level contribution between two vectors using cosine similarity.

Arguments:

- **factor**: 1D numpy array of shape `(n_timepoints,)` — independent variable
- **dependent**: 1D numpy array of shape `(n_timepoints,)` — dependent variable
- **mode**: Processing mode (1 = cosine similarity, 2 = angle/acos)

Returns:

- **contribution**: Scalar contribution value (float)

Listing 111: Trajectory-Level Contribution Calculation

```
contribution = tox_trajectory_contribution(factor, dependent, mode)
```

3. Spike Contribution Calculation

Computes timepoint-wise spike contributions between two vectors.

Arguments:

- **factor**: 1D numpy array of shape `(n_timepoints,)` — independent variable
- **dependent**: 1D numpy array of shape `(n_timepoints,)` — dependent variable

- **mode:** Processing mode (1 = Normal, 2 = RAP)

Returns:

- **contribution:** 1D numpy array of shape (n_timepoints,) with contribution values per timepoint (read-only)

Listing 112: Spike-Wise Contribution Calculation

```
spike_contributions = tox_spike_contribution(factor, dependent, mode)
```

4. Calculate Contributions (Allocation Version)

Calculates both spike and integrated contributions for a specific factor with internal memory allocation.

Arguments:

- **trajectories:** 3D numpy array of shape (n_factors, n_samples, n_timepoints) in Fortran order
- **i_factor:** 0-based index of the independent variable
- **dependent_idx:** 0-based index of the dependent variable
- **mode:** Processing mode (1 = Normal, 2 = RAP)

Returns:

- Dictionary containing:
 - **spikes:** Spike contributions [n_timepoints, n_samples] (read-only)
 - **trajectory:** Integrated contributions [n_samples] (read-only)

Listing 113: Contributions Calculation with Internal Allocation

```
result = tox_calc_contributions(trajectories, i_factor, dependent_idx,
                                mode)

# Access results
spike_contribs = result['spikes']           # [n_timepoints, n_samples]
traj_contribs = result['trajectory']        # [n_samples]
```

5. Calculate Contributions (Expert Version)

Calculates both spike and integrated contributions using caller-provided work arrays for optimal performance.

Arguments:

- **trajectories:** 3D numpy array of shape (n_factors, n_samples, n_timepoints) in Fortran order
- **i_factor:** 0-based index of the independent variable
- **dependent_idx:** 0-based index of the dependent variable
- **mode:** Processing mode (1 = Normal, 2 = RAP)
- **temp_factor_vector:** 1D work array of shape (n_timepoints,)
- **temp_dependent_vector:** 1D work array of shape (n_timepoints,)

Returns:

- Dictionary containing:
 - **spikes:** Spike contributions [n_timepoints, n_samples] (read-only)
 - **trajectory:** Integrated contributions [n_samples] (read-only)

Listing 114: Expert Contributions Calculation with Work Arrays

```
result = tox_calc_contributions_expert(
    trajectories, i_factor, dependent_idx, mode,
    temp_factor_vector, temp_dependent_vector)
```



```
)

# Access results (same structure as tox_calc_contributions)
spike_contribs = result['spikes']
traj_contribs = result['trajectory']
```

Processing Modes:

- mode = 1 (Normal): Returns cosine similarity for normal vectors
- mode = 2 (RAP): Returns arccos of cosine similarity for RAP projected vectors

Key Features of Python Interface:

- **Automatic Memory Management:** Handles Fortran memory layout and data type conversion automatically
- **Comprehensive Error Handling:** Converts Fortran error codes to Python exceptions
- **Memory Efficient:** Output arrays are marked as read-only to prevent accidental modification
- **Flexible Input:** Accepts both lists and numpy arrays with automatic conversion
- **Performance Optimized:** Expert versions allow pre-allocated work arrays for critical code paths

Usage Example:

Listing 115: Complete Contribution Analysis Workflow

```
import numpy as np

# Example data
n_factors, n_samples, n_timepoints = 5, 100, 50
trajectories = np.random.rand(n_factors, n_samples, n_timepoints).astype(
    np.float64)

# Single vector analysis
factor_vec = np.random.rand(n_timepoints)
dependent_vec = np.random.rand(n_timepoints)

# Compute trajectory contribution
traj_contrib = tox_trajectory_contribution(factor_vec, dependent_vec, mode
    =1)

# Compute spike contributions
spike_contribs = tox_spike_contribution(factor_vec, dependent_vec, mode=1)

# Full trajectory analysis
result = tox_calc_contributions(trajectories, i_factor=0, dependent_idx=1,
    mode=1)
spikes = result['spikes'] # [50, 100]
trajectory = result['trajectory'] # [100]

# Expert version with pre-allocated work arrays
temp_factor = np.empty(n_timepoints, dtype=np.float64)
temp_dependent = np.empty(n_timepoints, dtype=np.float64)
result_expert = tox_calc_contributions_expert(
    trajectories, i_factor=0, dependent_idx=1, mode=1,
    temp_factor, temp_dependent
)
```

1. Python Interface for Outlier Detection

The Python interface provides wrappers for the Fortran outlier detection routines, handling data type conversion, memory layout, and error checking automatically.

(a) Calculate Integrated Threshold (Expert Version)

Calculate empirical threshold for integrated contributions using pre-computed permutations for optimal performance.

Arguments:

- **contributions**: 1D numpy array of integrated contributions [n_samples] in Fortran order
- **percentile_val**: Percentile value for threshold (0.0–100.0)
- **permutation**: Pre-computed permutation indices for sorted contributions [n_samples] in Fortran order

Returns:

- **threshold**: Scalar threshold value

Listing 116: Integrated Threshold Calculation (Expert)

```
threshold = calc_integrated_threshold_expert(  
    contributions, percentile_val, permutation  
)
```

(b) Calculate Integrated Threshold (Convenience Version)

Calculate empirical threshold for integrated contributions with internal allocation and sorting.

Arguments:

- **contributions**: 1D numpy array of integrated contributions [n_samples] in Fortran order
- **percentile_val**: Percentile value for threshold (0.0–100.0)

Returns:

- **threshold**: Scalar threshold value

Listing 117: Integrated Threshold Calculation

```
threshold = calc_integrated_threshold(contributions, percentile_val)
```

(c) Detect Integrated Outliers

Identify outlier samples where integrated contributions exceed the empirical threshold.

Arguments:

- **contributions**: 1D numpy array of integrated contributions [n_samples] in Fortran order
- **threshold**: Scalar threshold value for outlier detection

Returns:

- **outlier_mask**: 1D boolean array indicating outlier samples [n_samples] (read-only)

Listing 118: Integrated Outlier Detection

```
outlier_mask = detect_outliers_integrated(contributions, threshold)
```

(d) Calculate Spike Thresholds (Expert Version)

Calculate empirical thresholds for spike contributions using pre-sorted permutations for optimal performance.

Arguments:

- **spike_contribs**: 2D numpy array of spike contributions [n_timepoints, n_samples] in Fortran order (rows=timepoints, columns=samples)
- **percentile_val**: Percentile value for threshold (0.0–100.0)
- **permutation**: Pre-computed permutation indices [n_samples, n_timepoints] in Fortran order (each column contains sorted indices for a timepoint)

Returns:

- **thresholds**: 1D numpy array of thresholds for each timepoint [n_timepoints] (read-only)

Listing 119: Spike Thresholds Calculation (Expert)

```
thresholds = calc_spike_thresholds_expert(
    spike_contribs, percentile_val, permutation
)
```

(e) **Calculate Spike Thresholds (Convenience Version)**

Calculate empirical thresholds for spike contributions with internal allocations and sorting.

Arguments:

- **spike_contribs**: 2D numpy array of spike contributions [n_timepoints, n_samples] in Fortran order (rows=timepoints, columns=samples)
- **percentile_val**: Percentile value for threshold (0.0–100.0)

Returns:

- **thresholds**: 1D numpy array of thresholds for each timepoint [n_timepoints] (read-only)

Listing 120: Spike Thresholds Calculation

```
thresholds = calc_spike_thresholds(spike_contribs, percentile_val)
```

(f) **Detect Spike Outliers**

Identify outliers in spike contributions by comparing against timepoint-specific thresholds.

Arguments:

- **spike_contribs**: 2D numpy array of spike contributions [n_timepoints, n_samples] in Fortran order (rows=timepoints, columns=samples)
- **thresholds**: 1D numpy array of thresholds for each timepoint [n_timepoints] in Fortran order

Returns:

- **outlier_mask**: 2D boolean array indicating outliers [n_timepoints, n_samples] (read-only)

Listing 121: Spike Outlier Detection

```
outlier_mask = detect_outliers_spike(spike_contribs, thresholds)
```

Key Features of the Python Interface:

- **Automatic Type Conversion**: Input arrays automatically converted to `np.float64` and `np.int32` with Fortran memory layout
- **Dimension Validation**: Input dimensions are validated before calling Fortran routines
- **Error Handling**: Fortran error codes are converted to Python exceptions via `check_err_code()`
- **Memory Efficiency**: Output arrays are marked as read-only to prevent accidental modification
- **Boolean Conversion**: C integer arrays (0/1) are automatically converted to Python boolean arrays

Usage Notes:

- Use **expert** versions when you have pre-computed permutations for better performance
- Use **convenience** versions for simpler workflows where internal allocation is acceptable
- All input arrays should be in **Fortran order** (column-major) for optimal performance
- Output arrays are **read-only** to maintain data integrity

2. Complete Pipeline Functions

The Python interface provides two complete pipeline functions for processing trajectories with integrated outlier detection workflows.

(a) **Process Trajectories with Per-Timepoint Percentiles**

Processes multiple trajectories with percentile thresholds calculated separately for each timepoint. This provides timepoint-specific sensitivity for spike contribution analysis.

Arguments:

- **trajectories**: 3D numpy array of shape (n_factors, n_samples, n_timepoints) in Fortran order
- **factor_mask**: 1D boolean array of shape (n_factors,) indicating which factors to process
- **dependent_idx**: 1-based index of the dependent variable (will be excluded from processing)
- **mode**: Processing mode (1 = Normal, 2 = RAP)
- **percentile**: Percentile value for threshold calculation (0.0–100.0)

Returns:

- Dictionary containing:
 - **integrated_contribs**: Integrated contributions [n_samples, n_processed_factors]
 - **spike_contribs**: Spike contributions [n_timepoints, n_samples, n_processed_factors]
 - **thresholds_integrated**: Thresholds for integrated contributions [n_processed_factors]
 - **thresholds_spike**: Thresholds for spike contributions [n_timepoints, n_processed_factors]
 - **outliers_integrated**: Outliers for integrated contributions [n_samples, n_processed_factors] (boolean)
 - **outliers_spike**: Outliers for spike contributions [n_timepoints, n_samples, n_processed_factors] (boolean)

Listing 122: Complete Trajectory Processing with Per-Timepoint Thresholds

```
results = tox_process_trajectories(  
    trajectories, factor_mask, dependent_idx, mode, percentile  
)  
  
# Access results  
integrated_contribs = results['integrated_contribs']      # [  
    n_samples, n_processed]  
spike_contribs = results['spike_contribs']                # [  
    n_timepoints, n_samples, n_processed]  
thresholds_integrated = results['thresholds_integrated']  # [  
    n_processed]  
thresholds_spike = results['thresholds_spike']           # [  
    n_timepoints, n_processed]  
outliers_integrated = results['outliers_integrated']      # [  
    n_samples, n_processed] (bool)  
outliers_spike = results['outliers_spike']               # [  
    n_timepoints, n_samples, n_processed] (bool)
```

(b) **Process Trajectories with Global Percentile**

Processes trajectories using a single global percentile threshold for all timepoints in spike contributions. This provides consistent sensitivity across all timepoints.

Arguments:

- **trajectories**: 3D numpy array of shape (n_factors, n_samples, n_timepoints) in Fortran order
- **factor_mask**: 1D boolean array of shape (n_factors,) indicating which factors to process
- **dependent_idx**: 1-based index of the dependent variable (will be excluded from processing)

- **mode**: Processing mode (1 = Normal, 2 = RAP)
- **percentile**: Percentile value for threshold calculation (0.0–100.0)

Returns:

- Dictionary containing:
 - **integrated_contribs**: Integrated contributions [n_samples, n_processed_factors]
 - **spike_contribs**: Spike contributions [n_timepoints, n_samples, n_processed_factors]
 - **thresholds_integrated**: Thresholds for integrated contributions [n_processed_factors]
 - **thresholds_spike**: Thresholds for spike contributions [n_processed_factors] (1D array)
 - **outliers_integrated**: Outliers for integrated contributions [n_samples, n_processed_factors] (boolean)
 - **outliers_spike**: Outliers for spike contributions [n_timepoints, n_samples, n_processed_factors] (boolean)

Listing 123: Complete Trajectory Processing with Global Threshold

```
results = tox_process_trajectories_flat(
    trajectories, factor_mask, dependent_idx, mode, percentile
)

# Access results - note thresholds_spike is 1D in flat version
integrated_contribs = results['integrated_contribs']      # [
    n_samples, n_processed]
spike_contribs = results['spike_contribs']                # [
    n_timepoints, n_samples, n_processed]
thresholds_integrated = results['thresholds_integrated']  # [
    n_processed]
thresholds_spike = results['thresholds_spike']            # [
    n_processed] (1D)
outliers_integrated = results['outliers_integrated']      # [
    n_samples, n_processed] (bool)
outliers_spike = results['outliers_spike']                # [
    n_timepoints, n_samples, n_processed] (bool)
```

Key Differences Between Pipeline Functions:

- **tox_process_trajectories**: Uses **per-timepoint percentiles** for spike thresholds (**thresholds_spike** shape: [n_timepoints, n_processed])
- **tox_process_trajectories_flat**: Uses **global percentile** for spike thresholds (**thresholds_spike** shape: [n_processed])

Pipeline Features:

- **Automatic Factor Filtering**: Processes only factors marked in **factor_mask**, excluding the dependent variable
- **Memory Efficient**: All output arrays use Fortran memory layout and are marked as read-only
- **Error Handling**: Comprehensive error checking with automatic exception conversion
- **Boolean Output**: Outlier masks are automatically converted from C integers to Python boolean arrays
- **Dimension Validation**: Input dimensions and factor mask consistency are validated

Usage Example:

Listing 124: Typical Pipeline Usage

```

import numpy as np

# Example data: 10 factors, 100 samples, 50 timepoints
trajectories = np.random.rand(10, 100, 50).astype(np.float64)
factor_mask = np.array([True, True, False, True, True, False, True, True,
                        , True, False])
dependent_idx = 5 # 1-based index
mode = 1 # Normal mode
percentile = 95.0

# Process with per-timepoint thresholds
results = tox_process_trajectories(
    trajectories, factor_mask, dependent_idx, mode, percentile
)

# Process with global threshold
results_flat = tox_process_trajectories_flat(
    trajectories, factor_mask, dependent_idx, mode, percentile
)

```

3. Python Core Contribution Analysis Functions

The Python interface provides comprehensive wrappers for trajectory contribution analysis, including EDF computation and various contribution calculation modes.

(a) Empirical Distribution Function (EDF) Computation

Computes the empirical cumulative distribution function for observed data values.

Arguments:

- **values:** Array of observed data values (list or numpy array)

Returns:

- Dictionary containing:
 - **unique_values:** Sorted unique data values (read-only numpy array)
 - **cdf_values:** Corresponding cumulative frequencies between 0 and 1 (read-only)
 - **n_unique:** Number of unique values found (int)

Listing 125: Empirical Distribution Function Computation

```

result = compute_edf(values)

# Access results
unique_vals = result['unique_values'] # Sorted unique values
cdf_vals = result['cdf_values']      # Cumulative distribution
values
n_unique = result['n_unique']        # Number of unique values

```

(b) Empirical Distribution Function (Expert Version)

Computes EDF using pre-sorted permutation for optimal performance.

Arguments:

- **values:** Array of observed data values (list or numpy array)
- **perm:** Pre-sorted permutation indices (1-based Fortran style, must be sorted by values[perm])

Returns:

- Dictionary containing:
 - **unique_values:** Sorted unique data values (read-only numpy array)

- `cdf_values`: Corresponding cumulative frequencies between 0 and 1 (read-only)
- `n_unique`: Number of unique values found (int)

Listing 126: Expert EDF Computation with Pre-sorted Permutation

```
result = compute_edf_expert(values, perm)

# Access results (same structure as compute_edf)
unique_vals = result['unique_values']
cdf_vals = result['cdf_values']
n_unique = result['n_unique']
```

3.2.7 Python Utils

1. Python BST Interface

(a) Build BST Index

Python wrapper for building BST index from NumPy arrays. Note that indices are 1-based and need to be converted for correct usage in python.

Input:

- `values` `numpy.ndarray`: 1D array of float64 values

Output:

- `np.array`: Permutation indices (1-based Fortran indices)

```
# - 1 at the end converts from fortrons 1-based indexing to pythons 0-based
indices = build_bst_index(values) - 1
```

(b) BST Range Query

Performs range queries on BST index. Accepts and returns Fortran 1-based indices.

Input:

- `values` `numpy.ndarray`: Original values array
- `indices` `numpy.ndarray`: BST index array from `build_bst_index` (1-based)
- `lower_bound` `float`: Lower bound of range (inclusive)
- `upper_bound` `float`: Upper bound of range (inclusive)

Output:

- Dictionary: (`matching_indices`, `count`) where indices are 1-based

```
# + 1 is added because of 0-based indexing. If the indices are already
# 1 based, this is not necessary
res = bst_range_query(values, indices + 1, lower_bound, upper_bound)
```

(c) Get Sorted Value

Note: This function is not implemented as a separate Python function. Use direct array indexing instead:

```
value = values[indices[position]] # position is 1-based
```

2. Python k-d Tree Interface

(a) Build k-d Tree Index

Python interface for constructing k-d trees. Requires Fortran-ordered arrays and returns 1-based indices.

Input:

- `points` `numpy.ndarray`: 2D array of points (dimensions $\times$ `n_points`) in Fortran order
- `dimension_order` `numpy.ndarray`: Dimension ordering (1-based Fortran indices, optional)

Output:

- `np.array`: Permutation indices (1-based Python indices)

```
kd_indices = build_kd_index(points, dimension_order)
```

(b) **Build Spherical k-d Tree**

Alias for `build_kd_index` with the same parameters and behavior.

Input: Same as `build_kd_index` **Output:** Same as `build_kd_index`

```
sphere_indices = build_spherical_kd(vectors, dimension_order)
```

(c) **Get Point from k-d Index**

Note: This function is not implemented as a separate Python function. Use direct array indexing instead:

```
point = points[:, kd_indices[position - 1]] # position is 1-based
```

3.3 R Pipeline

The R pipeline first reads the file into a raw, memory-efficient format (an integer array of ASCII codes) and then performs conversions from that raw format.

1. Read Gene IDs from TSV File (R)

The function reads gene identifiers from a tab-separated value (TSV) file, extracting specific gene IDs from a designated column while skipping header rows. This function is typically used as the first step in processing gene expression data to obtain the complete list of gene identifiers. Call the `read_gene_ids_from_tsv_file` function with the following arguments:

- `filename`: Character string of the filename
- `ngenes`: Number of gene IDs to read
- `gene_len`: Maximum length of each gene ID
- **(Optional)** `n_header_rows`: Number of header rows to skip
- **(Optional)** `gene_col`: Column index (1-based) of gene IDs in the file

Return list:

- `gene_ids`: Character vector of gene IDs read from the file
- `ierr`: Integer error code (0 if successful)

Listing 127: R Read Gene IDs from TSV File

```
read_gene_ids_from_tsv_file(filename, ngenes, gene_len, n_header_rows =
  1, gene_col = 1)
```

2. Read Expression Vectors (R)

The function reads gene expression data from multiple tabular files (CSV/TSV format), extracting expression values for specified genes across multiple samples. This function consolidates expression data from multiple files into a single matrix. Call the `read_expression_vectors` function with the following arguments:

- `file_list`: Character vector of filenames
- `gene_ids`: Character vector of gene IDs to match
- `n_header_rows`: Number of header rows to skip in each file
- `gene_col`: Column index (1-based) of gene IDs in each file

- **value_cols**: Vector of column indices (1-based) of expression values in each file
- **n_samples**: Total number of samples (rows) to read across all files
- **(Optional) delimiter**: Delimiter used in the files (default: tab)

Return list:

- **expression_vectors**: Numeric matrix of dimensions (n_samples × n_genes)
- **ierr**: Integer error code (0 if successful)

Listing 128: R Read Expression Vectors

```
read_expression_vectors(file_list, gene_ids, n_header_rows, gene_col,
  value_cols, n_samples, delimiter = "\t")
```

3. Read OrthoFinder File (R)

The function processes an OrthoFinder output file to map genes to their respective orthologous families. This function establishes the gene-family relationships essential for downstream orthology-based analyses. Call the `read_orthofinder_file` function with the following arguments:

- **filename**: Character string of the filename
- **gene_ids**: Character vector of gene IDs to match
- **n_families**: Number of unique gene families expected
- **family_len**: Maximum length of each family ID

Return list:

- **family_ids**: Character vector of family IDs read from the file
- **gene_to_fam**: Integer vector mapping each gene to its family index (0 = unassigned)
- **ierr**: Integer error code (0 if successful)

Note: If genes are not found in the gene IDs list a message will be printed but NO error code is thrown.

Listing 129: R Read OrthoFinder File

```
read_orthofinder_file(filename, gene_ids, n_families, family_len)
```

Complete R Data Loading Workflow Example:

Listing 130: R Complete Data Loading Workflow

```
# Step 1: Load gene identifiers
gene_result <- read_gene_ids_from_tsv_file(
  filename = "data/expression_data.tsv",
  ngenes = 50000,
  gene_len = 32,
  n_header_rows = 1,
  gene_col = 1
)
gene_ids <- gene_result$gene_ids

# Step 2: Load expression data from multiple samples
samples <- c("sample1.tsv", "sample2.tsv", "sample3.tsv", "sample4.tsv")
expression_result <- read_expression_vectors(
  file_list = samples,
  gene_ids = gene_ids,
  n_header_rows = 1,
  gene_col = 1,
  value_cols = 2,
```

```

    n_samples = 4
  )
  expression_vectors <- expression_result$expression_vectors

  # Step 3: Load orthology information
  family_result <- read_orthofinder_file(
    filename = "data/Orthogroups.tsv",
    gene_ids = gene_ids,
    n_families = 15000,
    family_len = 32
  )
  family_ids <- family_result$family_ids
  gene_to_fam <- family_result$gene_to_fam

```

4. Filter Unassigned Genes (R)

The function filters out genes that are not assigned to any family (where `gene_to_fam = 0`). This is essential for cleaning data before orthology-based analyses, ensuring only genes with valid family assignments are processed. Call the `filter_unassigned_genes` function with the following arguments:

- **gene_ids**: Character vector of gene IDs
- **expression_vectors**: Numeric matrix of expression values ($n_samples \times n_genes$)
- **gene_to_fam**: Integer vector mapping each gene to its family index (0 = unassigned)

Return list:

- **gene_ids**: Filtered character vector of gene IDs
- **expression_vectors**: Filtered numeric matrix of expression values
- **gene_to_fam**: Filtered integer vector mapping each gene to its family index
- **n_genes_kept**: Number of genes kept after filtering

Listing 131: R Filter Unassigned Genes

```
filter_unassigned_genes(gene_ids, expression_vectors, gene_to_fam)
```

5. Validate Gene to Family Mapping (R)

The function validates the gene-to-family mapping array, ensuring all family indices are within valid range and properly assigned. This prevents downstream errors from invalid family indices. Call the `validate_gene_to_family_mapping` function with the following arguments:

- **gene_to_fam**: Integer vector mapping each gene to its family index (0 = unassigned)
- **n_genes**: Number of genes
- **n_families**: Number of gene families

Listing 132: R Validate Gene to Family Mapping

```
validate_gene_to_family_mapping(gene_to_fam, n_genes, n_families)
```

6. Validate Expression Data (R)

The function validates expression data, checking for NaN/Inf values and optionally verifying non-negative values. This ensures data quality before computational analyses. Call the `validate_expression_data` function with the following arguments:

- **expression_vectors**: Numeric matrix of expression values ($n_samples \times n_genes$)
- **n_genes**: Number of genes
- **n_samples**: Number of samples

- **(Optional)** `check_non_negative`: Logical flag to check for non-negative values (default: TRUE)

Listing 133: R Validate Expression Data

```
validate_expression_data(expression_vectors, n_genes, n_samples, check_
  non_negative = TRUE)
```

7. Validate Family Centroids (R)

The function validates family centroids data, checking for NaN/Inf values and proper data structure. Centroids represent the central tendency of each gene family's expression profile. Call the `validate_family_centroids` function with the following arguments:

- `family_centroids`: Numeric matrix of family centroids ($n_samples \times n_families$)
- `n_families`: Number of gene families
- `n_samples`: Number of samples

Listing 134: R Validate Family Centroids

```
validate_family_centroids(family_centroids, n_families, n_samples)
```

8. Validate Shift Vectors (R)

The function validates shift vectors, ensuring data types are correct and structure matches expected patterns. Shift vectors represent the deviation of each gene from its family centroid. Call the `validate_shift_vectors` function with the following arguments:

- `shift_vectors`: Numeric matrix of shift vectors ($2 \times n_samples \times n_genes$)
- `expression_vectors`: Numeric matrix of expression values ($n_samples \times n_genes$)
- `family_centroids`: Numeric matrix of family centroids ($n_samples \times n_families$)
- `gene_to_fam`: Integer vector mapping each gene to its family index
- `n_samples`: Number of samples

Listing 135: R Validate Shift Vectors

```
validate_shift_vectors(shift_vectors, expression_vectors, family_
  centroids, gene_to_fam, n_samples)
```

9. Validate String Array Uniqueness (R)

The function validates that all strings in an array are unique, using a hashset internally. Note: This may increase memory usage temporarily for large datasets. Call the `validate_string_array_uniqueness` function with the following arguments:

- `string_arr`: Character vector of gene IDs
- `n_strings`: Number of genes

Listing 136: R Validate String Array Uniqueness

```
validate_string_array_uniqueness(string_arr, n_strings)
```

10. Validate All Data (R)

The function performs comprehensive validation of all data components in one call. This function executes all individual validations sequentially for complete data quality assurance. Call the `validate_all_data` function with the following arguments:

- `n_genes`: Number of genes
- `n_families`: Number of gene families

- `n_samples`: Number of samples
- `gene_ids`: Character vector of gene IDs
- `gene_family_ids`: Character vector of gene family IDs
- `gene_to_fam`: Integer vector mapping each gene to its family index (0 = unassigned)
- `expression_vectors`: Numeric matrix of expression values ($n_samples \times n_genes$)
- `family_centroids`: Numeric matrix of family centroids ($n_samples \times n_families$)
- `shift_vectors`: Numeric matrix of shift vectors ($2 \times n_samples \times n_genes$)

Listing 137: R Validate All Data

```
validate_all_data(n_genes, n_families, n_samples, gene_ids, gene_family_ids, gene_to_fam, expression_vectors, family_centroids, shift_vectors)
```

R Data Filtering and Validation Workflow Example:

Listing 138: R Data Filtering and Validation Workflow

```
# After loading data, filter unassigned genes
filter_result <- filter_unassigned_genes(gene_ids, expression_vectors, gene_to_fam)
gene_ids <- filter_result$gene_ids
expression_vectors <- filter_result$expression_vectors
gene_to_fam <- filter_result$gene_to_fam
n_genes_kept <- filter_result$n_genes_kept

# Run comprehensive validations
validate_string_array_uniqueness(gene_ids, n_genes_kept)
validate_string_array_uniqueness(family_ids, n_families)
validate_expression_data(expression_vectors, n_genes_kept, n_samples, check_non_negative = TRUE)
validate_gene_to_family_mapping(gene_to_fam, n_genes_kept, n_families)

# For complete analysis with centroids and shift vectors
validate_all_data(
  n_genes_kept, n_families, n_samples,
  gene_ids, family_ids, gene_to_fam,
  expression_vectors, family_centroids, shift_vectors
)

cat("All data validations passed successfully\n")
```

11. Create Zip Archive (Low-Level) (R)

The low-level function creates a zip archive from keys and filenames, directly calling the Fortran backend. Use this function for custom data serialization scenarios or when you need fine-grained control over archive creation. Call the `create_zip_archive` function with the following arguments:

- `zip_filename`: Name of the zip file to create
- `keys`: Vector of keys for the manifest (identifiers for each file)
- `filenames`: Vector of filenames to include in archive

Listing 139: R Create Zip Archive (Low-Level)

```
create_zip_archive(zip_filename, keys, filenames)
```

12. Extract Zip Archive (R)

The function extracts a zip archive created by `create_zip_archive` and returns a list mapping data keys to extracted filenames. This is useful for accessing individual files from custom archives. Call the `extract_zip_archive` function with the following arguments:

- `zip_filename`: Path to the zip file to extract

Returns:

- List mapping data keys to extracted temporary filenames

Listing 140: R Extract Zip Archive

```
extract_zip_archive(zip_filename)
```

13. Save TOX Data (High-Level) (R)

The high-level function saves TOX data to a zip archive, handling validation, serialization, and archive creation for standard-conform tox gene data. This function automatically manages temporary files and cleanup. Call the `save_tox_data` function with the following arguments:

- `zip_filename`: Name of the zip file to create
- **(Optional)** `gene_ids`: Character vector of gene IDs
- **(Optional)** `gene_ids_name`: Filename for gene IDs in the archive
- **(Optional)** `expression_vectors`: Numeric matrix of expression values ($n_samples \times n_genes$)
- **(Optional)** `expression_vectors_name`: Filename for expression vectors in the archive
- **(Optional)** `gene_to_fam`: Integer vector mapping each gene to its family index (0 = unsigned)
- **(Optional)** `gene_to_fam_name`: Filename for gene to family mapping in the archive
- **(Optional)** `family_ids`: Character vector of family IDs
- **(Optional)** `family_ids_name`: Filename for family IDs in the archive
- **(Optional)** `family_centroids`: Numeric matrix of family centroids ($n_samples \times n_families$)
- **(Optional)** `family_centroids_name`: Filename for family centroids in the archive
- **(Optional)** `shift_vectors`: Numeric matrix of shift vectors ($2*n_samples \times n_genes$)
- **(Optional)** `shift_vectors_name`: Filename for shift vectors in the archive

Listing 141: R Save TOX Data

```
save_tox_data(zip_filename,
              gene_ids = NULL, gene_ids_name = NULL,
              expression_vectors = NULL, expression_vectors_name = NULL,
              gene_to_fam = NULL, gene_to_fam_name = NULL,
              family_ids = NULL, family_ids_name = NULL,
              family_centroids = NULL, family_centroids_name = NULL,
              shift_vectors = NULL, shift_vectors_name = NULL)
```

14. Read TOX Data (R)

The function reads data from a zip archive created by `save_tox_data`, extracting and deserializing the requested data components for standard-conform tox gene data. This function automatically handles file extraction and cleanup. Call the `read_tox_data` function with the following arguments:

- `zip_filename`: Name of the zip file to read from
- **(Optional)** `gene_ids`: If not NULL, will attempt to read gene IDs from archive
- **(Optional)** `expression_vectors`: If not NULL, will attempt to read expression vectors from archive

- **(Optional) gene_to_fam:** If not NULL, will attempt to read gene to family mapping from archive
- **(Optional) family_ids:** If not NULL, will attempt to read family IDs from archive
- **(Optional) family_centroids:** If not NULL, will attempt to read family centroids from archive
- **(Optional) shift_vectors:** If not NULL, will attempt to read shift vectors from archive

Return list:

- **gene_ids:** Loaded gene IDs vector (if requested)
- **expression_vectors:** Loaded expression vectors matrix (if requested)
- **gene_to_fam:** Loaded gene-to-family mapping vector (if requested)
- **family_ids:** Loaded family IDs vector (if requested)
- **family_centroids:** Loaded family centroids matrix (if requested)
- **shift_vectors:** Loaded shift vectors matrix (if requested)

Listing 142: R Read TOX Data

```
read_tox_data(zip_filename,
              gene_ids = NULL,
              expression_vectors = NULL,
              gene_to_fam = NULL,
              family_ids = NULL,
              family_centroids = NULL,
              shift_vectors = NULL)
```

R Data Archiving Workflow Examples:

Listing 143: R Standard TOX Data Archiving

```
# Save complete TOX dataset
save_tox_data(
  zip_filename = "complete_tox_dataset.zip",
  gene_ids = gene_ids,
  gene_ids_name = "gene_ids.bin",
  expression_vectors = expression_vectors,
  expression_vectors_name = "expression.bin",
  gene_to_fam = gene_to_fam,
  gene_to_fam_name = "gene_to_fam.bin",
  family_ids = family_ids,
  family_ids_name = "family_ids.bin",
  family_centroids = family_centroids,
  family_centroids_name = "centroids.bin",
  shift_vectors = shift_vectors,
  shift_vectors_name = "shift_vectors.bin"
)

# Save partial dataset
save_tox_data(
  zip_filename = "basic_data.zip",
  gene_ids = gene_ids,
  gene_ids_name = "gene_ids.bin",
  expression_vectors = expression_vectors,
  expression_vectors_name = "expression.bin"
)
```

```
# Read data back with selective loading
loaded_data <- read_tox_data(
  zip_filename = "complete_tox_dataset.zip",
  gene_ids = TRUE,
  expression_vectors = TRUE,
  gene_to_fam = TRUE
)
```

Listing 144: R Custom Archive Creation and Extraction

```
# Create custom archive with non-standard data
create_zip_archive(
  zip_filename = "custom_data.zip",
  keys = c("3d_tensor", "results"),
  filenames = c("tensor_data.bin", "analysis_results.bin")
)

# Extract and access custom archive
file_mapping <- extract_zip_archive("custom_data.zip")
cat("Files in archive:\n")
print(file_mapping)
# Output: $3d_tensor = "tensor_data.bin", $results = "analysis_results.
bin"
```

Listing 145: R Mixed Standard and Custom Data Archiving

```
# Serialize custom arrays
tox_serialize_real_array(custom_array, "custom_data.bin")
tox_serialize_int_array(another_array, "another_data.bin")

# Create mixed archive with both standard and custom data
create_zip_archive(
  zip_filename = "mixed_archive.zip",
  keys = c("standard_expression", "custom_feature", "analysis_results"),
  filenames = c("expression.bin", "custom_data.bin", "another_data.bin")
)

# Later, extract and process
mapping <- extract_zip_archive("mixed_archive.zip")
expression_data <- tox_deserialize_real_array(mapping[["standard_
expression"]])
custom_features <- tox_deserialize_real_array(mapping[["custom_feature"
]])
results <- tox_deserialize_int_array(mapping[["analysis_results"]])
```

Important Notes for R Archive Functions:

- All archive functions automatically handle temporary file cleanup
- The `save_tox_data` function validates data dimensions before serialization
- Both `save_tox_data` and `read_tox_data` use the standard TOX data keys: "gene_ids", "expression", "gene_to_family", "family_ids", "family_centroids", "shift_vectors"
- Archives created with R functions are compatible with Python and Fortran implementations
- Error codes are automatically checked and converted to informative error messages

15. **TODO Functions relating section 1.2**

16. **Normalization of Expression Values**

The function runs a complete normalization pipeline on gene expression data. It first applies standard deviation normalization, then quantile normalization, followed by replicate averaging, and finally a log2 transformation. The result is returned in an array, which contains the fully normalized expression values. Call the `tox_normalization_pipeline` function with the following arguments:

- `input_matrix`: Numeric expression vector matrix with size `genes × tissues`
- `group_s`: Integer vector: start column index for each replicate group (1-based)
- `group_c`: Integer vector: number of columns per replicate group

Return:

- Numeric matrix: $\log_2(x+1)$ normalized expression

Listing 146: R Normalization Pipeline

```
tox_normalization_pipeline(input_matrix, group_s, group_c)
```

Alternatively you can calculate each step of the normalization pipeline manually by calling the following functions in sequence:

(a) **Normalize by Standard Deviation**

The function normalizes each gene's expression vector using `sqrt(mean(x^2))` across tissues (not classical standard deviation). Call the `tox_normalize_by_std_dev` function with the following arguments:

- `input_matrix`: A numeric matrix with genes as rows and tissues as columns

Returns:

- A normalized numeric matrix with the same dimensions and names as the input

Listing 147: R Standard Deviation Normalization

```
tox_normalize_by_std_dev(input_matrix)
```

(b) **Quantile Normalization**

The function performs quantile normalization of a gene expression matrix (F42-compliant) and computes average expression per rank across tissues. Call the `tox_quantile_normalization` function with the following arguments:

- `input_matrix`: A numeric matrix with genes as rows and tissues as columns

Returns:

- A quantile-normalized numeric matrix with the same dimensions and names as the input

Listing 148: R Quantile Normalization

```
tox_quantile_normalization(input_matrix)
```

(c) **Calculate Tissue Averages**

The function calculates tissue averages by averaging replicates within each group. For each group of tissue replicates, it computes the average expression per gene. The input matrix is column-major, flattened as a 1D array. Call the `tox_calculate_tissue_averages` function with the following arguments:

- `df`: A data frame or matrix with genes as rows and tissue replicates as columns

Returns:

- A data frame with genes as rows and averaged tissues as columns

Listing 149: R Tissue Averages Calculation

```
tox_calculate_tissue_averages(df)
```

(d) **Log2 Transformation**

The function applies $\log_2(x + 1)$ transformation to each element of the input matrix. This is computed via $\log(x + 1) / \log(2)$ for compatibility with WebAssembly (WASM). Call the `tox_log2_transformation` function with the following arguments:

- **input_matrix**: A numeric matrix with genes as rows and tissues as columns

Returns:

- A numeric matrix with log2-transformed expression values, preserving the same dimensions and names as the input.

Listing 150: R Log2 Transformation

```
tox_log2_transformation(input_matrix)
```

Additionally there are two helper functions to improve the data quality and prevent calculation errors:

(a) **Diagnose Data Quality**

The function checks your gene expression matrix for common data quality issues, such as missing values (NA), infinite values, NaNs, zeros, and negative values. It provides a summary of these problems, lists affected genes, and reports basic statistics like minimum, maximum, and mean values. Use this function to quickly assess the quality of your data before running normalization or analysis steps. Call the `tox_diagnose_data_quality` function with the following arguments:

- **input_matrix**: A numeric matrix with genes as rows and tissues as columns
- (Optional) **show_details**: Logical indicating whether to show detailed information (default is TRUE)

Returns:

- A list with diagnostic information about the data quality

Listing 151: R Diagnose Data Quality

```
tox_diagnose_data_quality(input_matrix, show_details)
```

(b) **Clean Data for Normalization**

The function prepares your gene expression matrix for normalization by handling problematic values. It can remove or impute NA, NaN, and Inf values, filter out genes with all zero values, and optionally convert very small values to zero. You can choose different strategies for dealing with missing data, ensuring your matrix is clean and ready for downstream analysis. Call the `tox_clean_data_for_normalization` function with the following arguments:

- **df_matrix**: A numeric matrix with genes as rows and tissues as columns
- (Optional) **remove_all_zero_genes**: Logical, whether to remove genes that are all zeros (default = TRUE)
- (Optional) **na_strategy**: Strategy for handling NA values: "remove_genes", "remove_samples", "impute_zero", "impute_mean" (default = "remove_genes")
- (Optional) **min_expression_threshold**: Minimum expression value to consider (values below this become 0 and default is 0)
- (Optional) **convert_small_to_zero**: Logical, whether very small numbers should be converted to zero (default = FALSE)

Returns:

- A cleaned matrix ready for normalization

Listing 152: R Clean Data for Normalization

```
tox_clean_data_for_normalization(df_matrix, remove_all_zero_genes, na_strategy, min_expression_threshold, convert_small_to_zero)
```

17. Calculate Fold Change

If fold change calculation is needed, this function can be used after normalization. The function calculates $\log_2$ fold changes between condition and control columns. For each control-condition pair, this function computes the $\log_2$ fold change by subtracting the expression value in the control column from the corresponding value in the condition column, for all genes. Call the `tox_calculate_fold_changes` function with the following arguments:

- `df`: A data frame with genes as rows and tissues/conditions as columns
- `control_pattern`: A string pattern to detect control columns
- `condition_patterns`: A character vector with patterns to detect condition columns

Returns:

- A data frame with genes as rows and $\log_2$ fold change values as columns

Listing 153: R Fold Change Calculation

```
tox_calculate_fc_by_patterns(df, control_pattern, condition_patterns)
```

18. Calculation of family Centroids

The function computes the centroid (mean expression vector) for each gene family in a gene expression dataset. It iterates over all families, selects the relevant genes for each family and calculates the mean vector for each family. There are two possible modes to use. To use all genes for calculation, select `all` mode. If you want to specify relevant genes for calculation, select `orthologs` mode. It is important that you provide a logical `ortholog_set` array when using `orthologs` mode. Call the `tox_group_centroid` function with the following arguments:

- `expression_vectors`: Expression vector matrix with size `n_axes` $\times$ `n_genes`
- `gene_to_family`: Gene-to-family mapping array with the length of `n_genes` (1-based indexing)
- `n_families`: Total number of gene families
- `mode`: Calculation mode ("all" or "orthologs")
- (Optional) `ortholog_set`: Logical array indicating if a gene is part of the calculation subset (only used in `orthologs` mode)

Return:

- `centroid_matrix`: Matrix of calculated family centroid vectors with size `n_axes` $\times$ `n_families`

Listing 154: R Calculation of Family Centroids

```
tox_group_centroid(expression_vectors, gene_to_family, n_families, mode, ortholog_set)
```

(a) Calculate Mean Vector

Alternatively you can calculate a single mean vector for a given set of vectors by calling the `mean_vector` function directly with the following arguments:

- `expression_vectors` `real(real64)`: Expression vector matrix with size `n_axes` $\times$ `n_genes`
- `gene_indices` `integer(int32)`: An array containing the column indices of the selected genes in `expression_vectors`

Return:

- **centroid:** Calculated mean vector (centroid) with length of **n_axes**

Listing 155: R Mean Vector Calculation

```
tox_mean_vector(expression_vectors, gene_indices)
```

19. Calculate Distance to Centroid

For each gene in the gene expression matrix, the function computes the Euclidean distance to its corresponding family centroid and returns a numeric vector of distances from each gene to its family centroid. Call the `tox_distance_to_centroid` function with the following arguments:

- **genes:** Gene expression matrix with size of (**n_genes** × **d**)
- **centroids:** Family centroid matrix with size of (**n_families** × **d**)
- **gene_to_fam:** Gene-to-family mapping integer vector with the length of **n_genes** (0 = no family, >0 = family index)
- **d:** Gene expression matrix dimension as integer

Listing 156: R Distance to Centroid Calculation

```
tox_distance_to_centroid(genes, centroids, gene_to_fam, d)
```

(a) Calculate Euclidean Distance

Alternatively you can calculate a single Euclidean distance between two vectors manually by calling `tox_euclidean_distance` function with the following arguments:

- **vec1:** First vector
- **vec2:** Second vector

Both vectors must have same length!

Listing 157: R Euclidean distance calculation

```
tox_euclidean_distance(vec1, vec2)
```

20. Detect Gene Outliers

The function first computes family scaling, then calculates relative distance indices (RDI) and identifies outliers based on the specified percentile threshold. It returns a list of components containing the outlier boolean values and intermediate results. Call the `tox_detect_outliers` function with the following arguments:

- **distances:** Numeric vector of distances for each gene to its centroid
- **gene_to_fam:** Gene-to-family mapping integer vector mapping with the same length as the distances vector (0 = no family, >0 = family index)
- **n_families:** Integer number of families
- **(Optional) percentile:** Percentile threshold value for outlier detection (default is 95.0)

Return list:

- **is_outlier:** Logical vector indicating whether each gene is an outlier
- **loess_x:** Numeric vector of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y:** Numeric vector of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **loess_n:** Integer vector with numbers of genes in each family used to compute the corresponding median and standard deviation

Listing 158: R Gene Outlier Calculation

```
tox_detect_outliers(distances, gene_to_fam, n_families, percentile)
```

Alternatively you can calculate each step of `tox_detect_outliers` manually with the following functions:

(a) Compute Family Scaling

The function calculates a scaling factor for each gene family by computing the median and standard deviation of gene-to-centroid distances within each family, then applies LOESS smoothing to these values. Call the `tox_compute_family_scaling` function with the following arguments:

- **distances**: Numeric vector of distances for each gene to its centroid
- **gene_to_fam**: Gene-to-family mapping integer vector mapping with the same length as the distances vector (0 = no family, >0 = family index)
- **n_families**: Integer number of families

Return dictionary:

- **dscale**: Numeric vector of scaling factors per family
- **loess_x**: Numeric vector of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y**: Numeric vector of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **indices_used**: Integer vector of indices of families that have been included in the calculation

Listing 159: R Compute Family Scaling

```
tox_compute_family_scaling(distances, gene_to_fam, n_families)
```

(b) Compute Family Scaling (Expert Version)

This function is the expert version of `tox_detect_outliers` and requires pre-allocated work arrays for maximum performance and control. It can be used when you need fine-grained control over memory allocation or you are calling this function many times in a tight loop. Call the `tox_compute_family_scaling_expert` function with the following arguments:

- **distances**: Numeric vector of distances for each gene to its centroid
- **gene_to_fam**: Gene-to-family mapping integer vector mapping with the same length as the distances vector (0 = no family, >0 = family index)
- **n_families**: Integer number of families
- **perm_tmp**: Permutation integer vector for sorting with the same length as distances
- **stack_left_tmp**: Stack integer vector for sorting with the same length as distances
- **stack_right_tmp**: Stack integer vector for sorting with the same length as distances
- **family_distances**: Work integer vector for sorting with the same length as distances

Return list:

- **dscale**: Numeric vector of scaling factors per family
- **loess_x**: Numeric vector of x coordinates used as reference points for LOESS smoothing (median distance)
- **loess_y**: Numeric vector of y coordinates used as reference points for LOESS smoothing (standard deviation)
- **indices_used**: Integer vector of indices of families that have been included in the calculation
- **perm_tmp**: Permutation vector for sorting

- **stack_left_tmp**: Stack vector for sorting
- **stack_right_tmp**: Stack vector for sorting
- **family_distances**: Work vector for sorting

Listing 160: R Compute Family Scaling (Expert Version)

```
tox_compute_family_scaling_expert(distances, gene_to_fam, n_families,
    perm_tmp, stack_left_tmp, stack_right_tmp, family_distances)
```

(c) **Compute Relative Distance Index (RDI)**

The function calculates the Relative Distance Index (RDI) for each gene by dividing its Euclidean distance to the family centroid by the scaling factor of its family. Call the `tox_compute_rdi` function with the following arguments:

- **distances**: Numeric vector of distances for each gene to its centroid
- **gene_to_fam**: Gene-to-family mapping integer vector mapping with the same length as the distances vector (0 = no family, >0 = family index)
- **dscale**: Numeric vector of scaling factors per family

Return list:

- **rdi**: Numeric vector containing the RDI value for each gene
- **sorted_rdi**: Numeric vector containing the RDI values sorted in ascending order

Listing 161: R Compute Relative Distance Index (RDI)

```
tox_compute_rdi(distances, gene_to_fam, dscale)
```

(d) **Identify Outliers**

The function identifies outliers based on the top percentile of RDI values. Call the `tox_identify_outliers` function with the following arguments:

- **rdi**: Numeric vector containing the RDI value for each gene
- **(Optional) percentile**: Percentile threshold number for outlier detection (default is 95.0)

Return list:

- **is_outlier**: Logical vector indicating whether each gene is an outlier
- **threshold**: Actual threshold value used for outlier detection

Listing 162: R Identify Outliers

```
tox_identify_outliers(rdi, percentile)
```

21. **Compute Normalized Tissue Versatility:**

The function calculates a normalized tissue versatility score for selected gene expression vectors. It is based on the angle between each gene expression vector and the space diagonal. Versatility is normalized to [0, 1], where 0 means uniform expression 1 means expression in only one axis. Call the `tox_compute_tissue_versatility` function with the following parameters:

- **expression_vectors**: Numeric matrix array with size `n_axes × n_vectors`
- **vector_selection**: Logical vector with length of `n_vectors` to specify which vectors should be included in the calculation
- **axes_selection**: Logical vector with length of `n_axes` to specify which axes should be included in the calculation

Return list:

- **tissue_versatilities**: Numeric vector containing the calculated tissue versatilities for selected vectors

- `tissue_angles_deg`: Numeric vector containing the calculated angles in degrees for selected vectors
- `n_selected_vectors`: Integer number of vectors processed
- `n_selected_axes`: Integer number of axes used in calculation

Listing 163: R Tissue Versatility Calculation

```
tox_calculate_tissue_versatility(expression_vectors, vector_selection,
                                axis_selection)
```

22. R BST Interface

(a) Build BST Index

R interface for building BST index. Uses R's native 1-based indexing.

Input:

- `x` numeric: Vector of real values

Output:

- `integer`: Permutation indices (1-based R indices)

```
ix <- build_bst_index(x)
```

(b) BST Range Query

Performs range queries on BST index. Returns a list with components.

Input:

- `x` numeric: Original values vector
- `ix` integer: BST index vector from `build_bst_index` (1-based)
- `lo` numeric: Lower bound of range (inclusive)
- `hi` numeric: Upper bound of range (inclusive)

Output:

- `list`: Contains indices and count components

```
result <- bst_range_query(x, ix, lo, hi)
indices <- result$indices # 1-based indices
count <- result$count
```

(c) Get Sorted Value

Retrieves value from sorted position using R 1-based indexing.

Input:

- `x` numeric: Original values vector
- `ix` integer: BST index vector (1-based)
- `position` integer: Position in sorted order (1-based)

Output:

- `numeric`: Value at specified position

```
value <- get_sorted_value(x, ix, position)
```

23. R k-d Tree Interface

(a) Build k-d Tree Index

R interface for k-d tree construction. Input matrix dimensions are rows=dimensions, columns=points.

Input:

- `X` matrix: Matrix of points (nrow = dimensions, ncol = n_points)
- `dim_order` integer: Dimension ordering (1-based, optional)

Output:

- **integer**: Permutation indices (1-based R indices)

```
kd_ix <- build_kd_index(X, dim_order)
```

(b) **Build Spherical k-d Tree**

R interface for spherical k-d tree construction with separate implementation.

Input:

- **V matrix**: Matrix of vectors (nrow = dimensions, ncol = n_vectors)
- **dim_order integer**: Dimension ordering (1-based, optional)

Output:

- **integer**: Permutation indices (1-based R indices)

```
sphere_ix <- build_spherical_kd(V, dim_order)
```

(c) **Get Point from k-d Index**

Retrieves point from k-d tree using R matrix indexing.

Input:

- **X matrix**: Original points matrix
- **kd_ix integer**: k-d tree index vector (1-based)
- **position integer**: Position in k-d tree order (1-based)

Output:

- **numeric**: Point values as vector

```
point <- get_kd_point(X, kd_ix, position)
```

24. R Interface for Array Serialization and Deserialization

The R interface provides bindings to the Fortran serialization and deserialization routines, enabling efficient handling of multi-dimensional arrays. This interface supports R arrays, matrices, and vectors with automatic type conversion to compatible Fortran data types.

(a) **General Architecture**

The R package provides ten main functions for array handling:

- **tox_serialize_int_array(array, filename)** - Serialize N-dimensional integer arrays
- **tox_deserialize_int_array(filename)** - Deserialize N-dimensional integer arrays
- **tox_serialize_real_array(array, filename)** - Serialize N-dimensional real arrays
- **tox_deserialize_real_array(filename)** - Deserialize N-dimensional real arrays
- **tox_serialize_char_array(array, filename)** - Serialize N-dimensional character arrays
- **tox_deserialize_char_array(filename)** - Deserialize N-dimensional character arrays
- **tox_serialize_logical_array(array, filename)** - Serialize N-dimensional logical array
- **tox_deserialize_logical_array(filename)** - Deserialize N-dimensional logical array
- **tox_serialize_complex_array(array, filename)** - Serialize N-dimensional complex array
- **tox_deserialize_complex_array(filename)** - Deserialize N-dimensional complex array

(b) **Array Functions**

Serialization Functions:

- **array (input)**: R array, matrix, or vector to be serialized

Type:

- **Integer**
- **Numeric** (double precision)

- Characters
- Logical
- Complex

Dimensions: Theoretically are N dimensions supported, but typical use cases should not exceed 5 dimensions

- **filename** (input): Name of the binary file to create
Type: character

Listing 164: R Array Serialization Example

```
tox_serialize_int_array(array, filename)
tox_serialize_real_array(array, filename)
tox_serialize_char_array(array, filename)
```

Deserialization Functions:

- **filename** (input): Name of the binary file to read
Type: character
- **Return Value:** Deserialized R array with original dimensions
Type: array, matrix, or vector
Data Type: Same as original

Listing 165: R Array Deserialization Example

```
restored_int <- tox_deserialize_int_array(filename)
restored_real <- tox_deserialize_real_array(filename)
restored_char <- tox_deserialize_char_array(filename)
```

(c) Implementation Details

- The R interface automatically converts between R data types and Fortran-compatible formats
- Integer vectors are converted to 32-bit integers
- Numeric vectors are converted to 64-bit floating point
- Character arrays handle variable-length strings with automatic padding
- All array dimensions are preserved during serialization/deserialization
- The interface handles both regular arrays and matrices seamlessly

Usage Example

Listing 166: Full example of serialization and deserialization in R

```
tox_serialize_int_array(arr1, fn("int1d.bin"))
restored <- tox_deserialize_int_array(fn("int1d.bin"))

# Real array example
arr2 <- matrix(runif(12), nrow = 3, ncol = 4)
tox_serialize_real_array(arr2, fn("real2d.bin"))
restored <- tox_deserialize_real_array(fn("real2d.bin"))

# Character array example
arr3 <- c("GENE1", "BRCA1", "TP53", "EGFR")
tox_serialize_char_array(arr3, fn("char1d.bin"))
restored <- tox_deserialize_char_array(fn("char1d.bin"))
```


(d) **Compatibility Notes**

- Requires R with appropriate Fortran bindings
- Supports standard R data types (integer, numeric, character, logical, complex)
- Preserves array attributes and dimensions during serialization
- Character arrays maintain content but may have padding removed upon deserialization

3.4 Fortran utilities

1. Hashing

The TensorOmics package provides high-performance hashing utilities using the XXH3 algorithm for both hash maps and hash sets. These utilities are implemented in Fortran with separate chaining collision resolution and automatic resizing.

(a) **Hashmap**

The hashmap implementation provides key-value storage with string keys and integer values, using the XXH3 hashing algorithm for optimal performance.

- **Type Definition:**

```
type(hashmap_type) :: map
```

- **Create Hashmap:** Initialize a new hashmap with optional initial size. The Hashmap size will be set to the next greater power of two.

```
call hashmap_create(map, initial_size)
```

- map: Hashmap object to create
- **(Optional) initial_size:** Initial size of the hashmap

- **Insert Key-Value Pair:** Insert or update a key-value pair in the hashmap.

```
call hashmap_put(map, key, value)
```

- map: Hashmap to insert into
- key: String key to store
- value: Integer value to store

- **Lookup Key:** Retrieve the value associated with a key.

```
value = hashmap_get(map, key)
```

- map: Hashmap to search
- key: String key to look for
- **Returns:** Integer value, or -1 if key not found

- **Destroy Hashmap:** Clean up and deallocate the hashmap. It is important to always call this function when the hashmap is no longer in use to avoid memory leaks.

```
call hashmap_destroy(map)
```

- map: Hashmap to destroy

Hashmap Usage Example:

```
use xxh3_hashmap_module

type(hashmap_type) :: gene_map
integer :: gene_index
```

```

! Create hashmap for gene IDs
call hashmap_create(gene_map, initial_size=1000)

! Store gene indices
call hashmap_put(gene_map, "Gene_001", 1)
call hashmap_put(gene_map, "Gene_002", 2)
call hashmap_put(gene_map, "Gene_003", 3)

! Lookup gene index
gene_index = hashmap_get(gene_map, "Gene_002")
! gene_index = 2

! Clean up
call hashmap_destroy(gene_map)

```

(b) Hashset

The hashset implementation provides unique string storage with fast membership testing, using the same XXH3 hashing algorithm as the hashmap.

- **Type Definition:**

```
type(hashset_type) :: set
```

- **Create Hashset:** Initialize a new hashset with optional initial size. The hashset size will be increased to the next greatest power of two.

```
call hashset_create(set, initial_size)
```

- `set`: Hashset object to create
- **(Optional)** `initial_size`: Initial size of the hashset

- **Insert Key:** Insert a unique key into the hashset.

```
call hashset_put(set, key, ierr)
```

- `set`: Hashset to insert into
- `key`: String key to store
- `ierr`: Error code (0 if successful, error if duplicate)

- **Check Membership:** Test if a key exists in the hashset.

```
exists = is_in_hashset(set, key)
```

- `set`: Hashset to search
- `key`: String key to check
- **Returns:** Logical indicating membership

- **Destroy Hashset:** Clean up and deallocate the hashset. It is important that this function is always called when the hashset is no longer in use to avoid memory leaks.

```
call hashset_destroy(set)
```

- `set`: Hashset to destroy

Hashset Usage Example:

```

use xxh3_hashmap_module

type(hashset_type) :: unique_genes
integer :: ierr

```

```

logical :: exists

! Create hashset for unique gene validation
call hashset_create(unique_genes, initial_size=5000)

! Add genes to set
call hashset_put(unique_genes, "Gene_001", ierr)
call hashset_put(unique_genes, "Gene_002", ierr)
call hashset_put(unique_genes, "Gene_003", ierr)

! Check for duplicates (will return error)
call hashset_put(unique_genes, "Gene_001", ierr)
! ierr != 0 indicates duplicate key

! Test membership
exists = is_in_hashset(unique_genes, "Gene_002")
! exists = .true.

! Clean up
call hashset_destroy(unique_genes)

```

Technical Features:

- Uses XXH3 non-cryptographic hash algorithm for optimal performance
- Automatic resizing when load factor exceeds 0.75
- Separate chaining collision resolution
- Table sizes are powers of two for efficient modulo operations
- Keys are automatically trimmed and normalized
- Memory efficient with proper cleanup procedures
- Thread-safe for read operations (concurrent writes not supported)

Performance Characteristics:

- Average case $O(1)$ for insertions and lookups
- Memory usage proportional to number of elements + overhead
- Resizing operations are $O(n)$ but amortized over many insertions

2. Sorting Algorithms

The TensorOmics package provides high-performance quicksort implementations for various data types using iterative algorithms with manual stack management. These sorting routines operate indirectly through permutation vectors to preserve the original data order.

(a) Generic Sorting Interface

A unified interface for sorting different data types through the `sort_array` generic procedure.

• Generic Procedure:

```
call sort_array(array, perm, stack_left, stack_right)
```

- `array`: Input array to sort (real, integer, or character)
- `perm`: Permutation vector that will be reordered
- `stack_left`: Work array for left indices (size = array size)
- `stack_right`: Work array for right indices (size = array size)

Supported Types:

```

! Real arrays
real(real64) :: data_real(:)
! Integer arrays

```

```
integer(int32) :: data_int(:)
! Character arrays
character(len=*) :: data_char(:)
```

(b) Type-Specific Sorting Procedures

Direct access to type-specific sorting routines for explicit control.

- **Sort Real Arrays:**

```
call sort_real(array, perm, stack_left, stack_right)
```

- array: Real(real64) input array to sort
- perm: Permutation vector that will be sorted
- stack_left: Manual stack of left indices
- stack_right: Manual stack of right indices

- **Sort Integer Arrays:**

```
call sort_integer(array, perm, stack_left, stack_right)
```

- array: Integer(int32) input array to sort
- perm: Permutation vector that will be sorted
- stack_left: Manual stack of left indices
- stack_right: Manual stack of right indices

- **Sort Character Arrays:**

```
call sort_character(array, perm, stack_left, stack_right)
```

- array: Character input array to sort (lexicographic order)
- perm: Permutation vector that will be sorted
- stack_left: Manual stack of left indices
- stack_right: Manual stack of right indices

Sorting Usage Example:

```
use f42_utils
use iso_fortran_env, only: real64, int32

real(real64) :: expression_data(1000)
integer(int32) :: perm(1000), stack_left(1000), stack_right(1000)
integer :: i

! Initialize permutation
do i = 1, 1000
    perm(i) = i
end do

! Sort expression data indirectly
call sort_real(expression_data, perm, stack_left, stack_right)

! Access sorted data through permutation
! expression_data(perm(1)) is the smallest value
! expression_data(perm(1000)) is the largest value
```

(c) Technical Implementation Details

- **Algorithm:** Iterative quicksort with manual stack management
- **Sorting Method:** Indirect sorting via permutation vectors

- **Pivot Selection:** Median-of-three (center element)
- **Memory:** Requires two work arrays of same size as input
- **Performance:** Average case $O(n \log n)$, worst case $O(n^2)$
- **Stability:** Not stable (order of equal elements not preserved)

(d) **Work Array Requirements**

For sorting an array of size n , you must provide:

- `perm(n)`: Permutation vector (initialize with 1, 2, ..., n)
- `stack_left(n)`: Work array for stack operations
- `stack_right(n)`: Work array for stack operations

Initialization Pattern:

```
do i = 1, n
    perm(i) = i           ! Identity permutation
end do
```

(e) **Advanced Usage: Direct Algorithm Access**

For maximum performance in tight loops, you can call the internal quicksort procedures directly:

```
! Internal quicksort for real arrays
call quicksort_real(array, perm, n, stack_left, stack_right)

! Internal quicksort for integer arrays
call quicksort_int(array, perm, n, stack_left, stack_right)

! Internal quicksort for character arrays
call quicksort_char(array, perm, n, stack_left, stack_right)
```

- These procedures require explicit array size parameter n
- Work arrays must be properly initialized
- Provides slightly better performance by skipping interface overhead

Performance Considerations:

- Manual stack management avoids recursion limits for large arrays
- Indirect sorting preserves original data and allows multiple sort orders
- Character sorting uses Fortran's intrinsic lexicographic comparison
- Memory usage is $4n$ for all data including input, work array, and both stacks.
- In case of memory considerations, heapsort might be used.

Error Handling:

- All sorting procedures are **pure** (no side effects)
- No explicit error codes - assumes valid input dimensions
- User must ensure work arrays are properly sized
- Invalid permutations may result in incorrect sorting